

Servey - SEO Context

The SEO playbook for servey.in: what we do, what we deliberately do not do, and the evidence behind each decision.

Companion to `CONTEXT.md`, which covers the product and codebase. This file covers search and answer-engine strategy only.

- **Last reviewed:** 1 September 2026
- **Source of truth:** `docs/SEO-CONTEXT.md` . `docs/seo-context.pdf` is generated from it by `npm run context:pdf` - never edit the PDF.

Why this file exists. Nearly every conclusion below was expensive to reach and easy to forget. Several were reached *twice* because the reasoning was not written down. If you are about to "improve" something here, read the reason first - most of these are deliberate trade-offs, not oversights.

1. Baseline - where we actually are

Google Search Console, 3-month window to **29 August 2026** (read 31 Aug):

Metric	Value	25 Aug	14 Aug
Clicks	54	38	16
Impressions	2,200	1,460	626
CTR	2.5%	2.6%	2.6%
Average position	16.5	16.4	15.8
Distinct queries	162	39	-

Clicks +42% and impressions +51% in four days; the query count more than quadrupled. The query count is the number that matters most here. Going from 39 to 162 distinct queries means Google has stopped treating us as a handful of pages about Jump Desktop and started matching us across the whole category.

Position holding flat at ~16.5 while impressions grow is the expected shape and is **not** a regression - see the note in §10 on reading position correctly.

Brand vs non-brand: `servey` is 9 clicks / 257 impressions - **17% of clicks from 12% of impressions.**

Brand share of clicks has gone 44% -> 24% -> 17% over three readings while brand click volume stayed flat at 9. Non-brand is now carrying essentially all growth. That was the goal of this whole strategy.

Attribution caveat: the 162-row query table accounts for only **23 of the 54 clicks**. GSC anonymises rare queries, so ~31 clicks are unattributable to a specific term. Cluster click counts below are therefore floors, not totals.

Index coverage (verified 31 Aug 2026):

- **Google** - sitemap submitted, now **34 URLs**.
- **Bing** - indexed; ChatGPT Search reads Bing's index, so that pipe is connected.

The core diagnosis - CORRECTED 1 Sep 2026

The site is bimodal, and the 16.5 average is a meaningless midpoint between two populations that behave completely differently. Per-query position data pulled 1 Sep (the metric was never enabled before, which is how this went unnoticed for three readings):

Query	Position	Impressions	Clicks
splashtop vs jump desktop	2.3	9	1
jump desktop vs rustdesk	3.3	25	1
jump desktop vs chrome remote desktop	3.3	19	1
rustdesk vs jump desktop	3.8	17	4
jump desktop vs splashtop	4.0	8	1
servey (brand)	5.0	257	9
screens vs jump desktop	7.0	27	3
jump desktop vs screens 5	7.9	24	1
screens 5 vs jump desktop	10.8	24	1
remote control iphone from mac	44.6	23	0
remote access mac from iphone	40.0	18	0

The comparison cluster is already on page one - top five for most of it. The core product cluster is on page four to five. Averaging those two gives 16.5, a number that describes nothing real and that we have been reasoning from since 14 August.

What this changes:

1. **"We are stuck on page two, it is an authority problem" was wrong as stated.** We are on page one where we have topical depth (six Jump Desktop comparisons) and on page four where we have one page. Authority is still the binding constraint on the core cluster, but the comparison cluster proves the site can rank when the topic is covered properly.
2. **The §4 ruling against metadata work was based on the bad average and is now partially lifted** - see that section. CTR at position 3 is not evaluated against the same benchmark as CTR at position 16.
3. **The core-cluster rebuild of 31 Aug cannot be judged on CTR.** At position 40 the page is not being seen, so a flat click count over the next fortnight means nothing. Watch **position**, not clicks, for that page.

Index coverage (verified 1 Sep 2026): 28 indexed, 3 not indexed (2 "Page with redirect", 1 "Discovered - not indexed"). **"Crawled - currently not indexed" is 0**, which is the bucket that carries a quality verdict. There is **no index bloat** and nothing has been rejected on quality grounds.

2. What the query data proves

All 162 queries, clustered (3-month window to 29 Aug 2026):

Cluster	Phrasings	Impressions	Clicks	Verdict
Competitor vs competitor (Jump Desktop is in nearly all of them)	~20	~235	14	✅ Earns every attributable non-brand click
Control Mac FROM iPhone (remote access mac from iphone , control mac with iphone , ...)	55	~120	0	❌ Our core product intent. Biggest miss on the site
Control iPhone FROM Mac (remote control iphone from mac , ...)	13	~46	0	❌ Inverse intent - we had no page for it at all
Headless Mac mini (headless mac mini setup , macos headless mode , ...)	12	~32	0	❌ How-to intent; we only had a "why" page
AI on Mac (mac ai agent , run local ai on mac)	10	~26	0	Three posts already; saturated, do not add more
Screens 5 pricing (screens 5 cost , screens 5 price)	4	~17	0	Commercial intent, uncovered. Blocked: pricing not verifiable
airplayuiagent , publicrelay , splashbuddy	5	~8	0	Accidental - macOS process names. Ignore

The finding that reset priorities on 31 Aug: the "control my Mac from my iPhone" family - Servey's exact job - produces ~120 impressions across 55 distinct phrasings and converts zero of them. It is by far the largest cluster on the site and the only one where we should be the definitive answer.

The cause was diagnosable by reading the page. `control-your-mac-from-iphone-ipad` was written as a *pitch* ("what good Mac remote control should feel like"), not as a *guide*. It never mentioned System Settings, Screen Sharing, VNC or any actual step. Google ranked it for the topic and then buried it, because the queries are overwhelmingly `how to ...` and the page answered nothing. Rebuilt 31 Aug as a real how-to that leads with Apple's free built-in route.

Four conclusions that drive every content decision:

1. **Comparison content is the only proven format** - specifically *competitor-vs-competitor*, where Servey is **not** the subject. We rank by genuinely helping someone choose between two other products.
2. **Fragmentation is the failure mode.** 55 phrasings of one intent with zero clicks is what Google does when it reads a site as marginally relevant to all of a group and authoritative on none.
3. **A page that pitches instead of answering cannot rank for a how-to query**, however well optimised. Lead with the free/built-in answer, then position Servey against the part it genuinely does not solve. This is exactly the shape that made `does-mac-screen-sharing-work-over-the-internet` convert.
4. **Check for inverse intent.** ~46 impressions were people wanting the *opposite* direction (their iPhone from their Mac). Serving that on its own page both captures the traffic and stops it diluting the core page.

3. What we DO

Content

- **Write competitor-vs-competitor comparisons.** The proven format. Currently: `screens-vs-jump-desktop`, `jump-desktop-vs-teamviewer`, `jump-desktop-vs-rustdesk`, `rustdesk-vs-anydesk`, `screens-5-alternatives`, `screens-jump-desktop-alternative-mac`, `termius-alternative-mac-terminal`, `best-remote-desktop-for-mac`.
- **Lead comparison posts with a `table` block.** Tables are the structure answer engines extract most reliably.
- **State where every competitor is genuinely better.** This is the single reason these pages get read and cited. It is a strategy, not politeness.
- **Keep Servey out of the hero slot** in competitor-vs-competitor posts: one late mention, labelled pre-launch, ending by telling the reader to buy an established product if they need one today.
- **One commercial pillar per intent**, with informational blog posts as spokes pointing at it (`/control-mac-from-iphone` ← `control-your-mac-from-iphone-ipad`).
- **Name Servey in the first ~15% of the article, and measure it.** Two metrics, both easy to check by stripping tags and counting character offsets: **first-mention depth** $\leq 15\%$ and **density 4.0-6.0 per 1k words** (a short post will run a little over on density; depth is the one that matters). The informational posts both drifted to 56% and 63% because they had no reason to name Servey until the dedicated section - a reader on a phone scrolls a long way before learning a product exists. The fix that works is **structural, not a pitch**: give the post's table a Servey column, or name Servey inside the row or list item it genuinely belongs to, alongside the competitors.
- **On pages used for outreach, keep the honest "No".** The Back to My Mac table says Servey does *not* mount your Mac's disk in the Finder, because it doesn't. A table that only says yes is worth nothing to a reader, and an editor deciding whether to swap a dead citation for our page will not link to an advert.

On-page

- `metaTitle` ≤ 51 chars (the `%s - Servey` template adds 9; rendered ≤ 60).
- Descriptions **110–160 chars**.
- Explicit `alternates.canonical` **and** explicit `openGraph.url` on every page.
- One `<h1>`, no skipped heading levels.
- Shared `ogImage` object (never the bare string, which drops width/height/alt).

Technical

- **Static generation everywhere.** Every route is Static/SSG; guard this.
- **Speculation Rules** - same-origin prefetch/prerender at `moderate` eagerness, excluding `/api/` and `/ingest/`.
- **Structured data:** one `@graph` (Organization → WebSite → SoftwareApplication with PreOrder Offers → FAQPage → ImageObject), plus per-page `Article` / `WebPage` / `BreadcrumbList` / `FAQPage`. Cross-reference by `@id`; never re-declare a node inline.
- **sizes** on every responsive image, matched to its real rendered width.

- **IndexNow** after every content deploy (`npm run indexnow`) - fastest route into Bing, and therefore into ChatGPT Search.
- `npm run seo:audit` before/after deploys. Currently **32/32 pass**.

Answer engines (AEO / GEO)

- **Self-contained passages of ~130–170 words** that still make sense quoted in isolation.
 - **Question-shaped H2s** mirroring real queries.
 - **Direct answer in the first sentence** after each heading, then support.
 - `robots.ts` **explicitly allows AI crawlers by name** - GPTBot, OAI-SearchBot, ChatGPT-User, ClaudeBot, Claude-Web, Claude-SearchBot, PerplexityBot, Perplexity-User, Google-Extended, Applebot-Extended, Bingbot, Amazonbot, CCBot.
-

4. What we deliberately DO NOT do

This section is as important as the last one. Each of these is a considered decision with a reason.

Never - integrity

We don't	Why
Invent ratings, testimonials, review counts, user numbers, or <code>AggregateRating</code>	Survey is pre-launch with zero public users. <code>AggregateRating</code> without real reviews is a Google structured-data violation, and the rest is simply false.
Fabricate product screenshots	Every image is a real capture or an honestly-styled placeholder.
Publish a comparison that hides where a competitor wins	It is the entire credibility asset. A post that reads like an advert has failed at its job.
Imply Termius shipped malware	They were the victim of the ZuRu impersonation campaign. Saying otherwise is false and defamatory.
Solicit upvotes anywhere	Against Product Hunt / AlternativeTo rules; triggers rank penalties.
Add our own links to Wikipedia (dead-link building)	Verified 24 Aug: Wikipedia renders external links as <code>rel="nofollow"</code> sitewide, so they pass no ranking signal. Adding our own also violates WP:COI/WP:REFSPAM, where the failure mode is <code>servey.in</code> on the public Wikipedia spam blacklist. Wikipedia is downstream of independent press, not a substitute. Broken-link building off Wikipedia is fine and is our active tactic - see <code>docs/BACKLINKS.md</code> .
Buy links, use PBNs, or exchange links	Cheap short-term, fatal long-term.
Work a bulk "N places to get backlinks" list top to bottom	Triaged one of 75 on 26 Aug: 9 were usable now . The rest were editorial placements you must earn (Forbes, TechCrunch), revenue-gated directories we cannot enter pre-launch, review sites needing real customers, open-source-only listings, or content farms. Submitting to dozens of unrelated low-quality directories in a burst is the classic footprint of a link scheme, and on a domain with 2 referring domains it is the <i>only</i> pattern in our profile. Pick for relevance, submit steadily, and never post the same text everywhere without a canonical. See <code>docs/BACKLINKS.md</code> §6.
Keyword-stuff, cloak, or hide text	Trivially detected; the fragmentation fix is pillar pages, not stuffing.
Publish thin AI-generated filler to inflate page count	Google's QRG §4.6.5–6 targets exactly this. Page count is not the goal.

Not now - strategy

We don't	Why
Chase "recommend me a tool" AI queries	Verified 16 Aug: ChatGPT answered "tool to screen share to my Mac from my phone" with Chrome Remote Desktop, AnyDesk, TeamViewer. Correctly - Servey has nothing to download, so recommending it would be a worse answer. Structurally unwinnable until launch. Do not treat it as an SEO failure.
Chase "termius alternative" broadly	It would pull multi-host-SSH searchers who bounce, polluting waitlist quality pre-launch. The honest overlap is narrow: "run a terminal on your own Mac without SSH setup".
Compete on Termius's page axis (SSH client x platform)	They hold #1 there. We copied their <i>template technique</i> onto a different axis - task x device - where they do not rank at all.
Add more pages targeting "control Mac from iPhone"	That recreates the exact fragmentation <code>/control-mac-from-iphone</code> was built to fix. One pillar, many spokes.
Rewrite metadata sitewide to chase CTR	Still ruled out as a <i>bulk</i> exercise. But the original reasoning ("above par for position 16.4") used the misleading site average and is void - see §1. Targeted work on the specific page-one comparison URLs is now allowed, and only those.
Build doorway pages per query variant	The 10 phrasings get absorbed by one strong page, not 10 weak ones.
Do multilingual / translated SEO	Considered 24 Aug and deliberately held. 19 posts x 7 locales is 133 URLs on a domain with 2 referring domains already averaging position 16.3. Translations multiply pages, not authority, and scaled machine translation without human review is named in Google's spam policies. The tactic works for language-independent utility queries ("instagram photo downloads"); ours are niche Mac comparison queries whose product names, docs and buyers are all English. Revisit once referring domains are in double digits, then as a one-language, 4-post pilot chosen from GSC Countries data - never a 7-locale bulk translation.
Chase <code>airplayuiagent</code> traffic	Accidental, off-intent, and converts nobody.

Myths we do not act on

Belief	Reality
llms.txt improves AI citation	It does not. No major engine consumes it; this is evidence-based, not opinion. We keep ours because it costs nothing and is a tidy human-readable map - but it is not a citation lever and must not be cited as one.
FAQ schema wins SERP real estate	Retired 7 May 2026. See §7.
More pages = more traffic	39 queries across 26 pages says relevance is fine; authority is the constraint.
Content needs chunking / AI-specific keyword rewriting	Engines handle synonyms; this is wasted effort.

5. Cluster → page ownership map

Each intent has exactly one owner. Adding a second competitor for any row is how fragmentation comes back.

Intent cluster	Owning page	Supporting spokes
Control Mac from iPhone	/control-mac-from-iphone	blog/control-your-mac-from-iphone-ipad (the how-to half - see §2)
Control iPhone from Mac (inverse intent)	blog/control-iphone-from-mac	-
Terminal from iPhone	/terminal-on-iphone	blog/real-terminal-on-your-mac-from-iphone
Terminal from iPad	/terminal-on-ipad	-
Headless Mac mini / macOS headless mode	/headless-mac-mini	blog/headless-mac-mini-setup (setup), blog/control-a-headless-mac-mini-remotely (access)
Run local AI on Mac / mac ai agent	/remote-mac-for-ai-agents	blog/run-ai-agents-locally-on-your-mac , blog/run-ai-agents-on-your-mac-remotely , blog/stay-in-control-of-ai-agents-from-anywhere
Developer remote Mac access	/mac-for-developers	blog/who-is-survey-for-developers-home-labs
Mac home lab	/mac-home-lab	-
Cellular / strict networks	blog/access-your-mac-remotely-over-cellular	-
Screens vs Jump Desktop	blog/screens-vs-jump-desktop	blog/screens-jump-desktop-alternative-mac
Back to My Mac replacement	blog/what-replaced-back-to-my-mac	blog/does-mac-screen-sharing-work-over-the-internet
RustDesk vs AnyDesk	blog/rustdesk-vs-anydesk	-
Screens 5 alternatives	blog/screens-5-alternatives	-
Jump Desktop vs TeamViewer	blog/jump-desktop-vs-teamviewer	-
Jump Desktop vs RustDesk	blog/jump-desktop-vs-rustdesk	-
Termius / SSH overlap	blog/termius-alternative-mac-terminal	-
Category roundup	blog/best-remote-desktop-for-mac	-

6. Internal linking - a quirk that matters

`p` blocks render as plain text. `app/blog/[slug]/page.tsx` parses no markdown and no HTML, so **you cannot put an anchor in body copy**. The only real internal-link mechanism between posts is **exact-string keyword overlap feeding** `relatedPosts()` .

Engineering the `keywords` array *is* engineering the link graph. When adding a post, deliberately share keywords with the posts you want linking to it. The three newest comparison posts were given keyword overlap with `screens-vs-jump-desktop` specifically because it is the one page earning clicks, so its authority flows to them.

Use-case pages *do* have real anchors: the cross-link mesh is derived automatically from the `useCases` array order, so adding a page wires it in.

7. Structured-data deprecation register

Type	Status	Action
FAQPage	Rich results retired 7 May 2026. SC report + Rich Results Test support removed June 2026; API data August 2026.	Keep the markup. It remains a valid Schema.org type, Google confirms unused structured data does not affect Search, and answer engines still read it. Only the <i>expectation</i> changed.
HowTo	Removed June 2025	Do not add
SpecialAnnouncement , ClaimReview , VehicleListing , EstimatedSalary , LearningVideo , Course carousel, Book Actions	Removed June 2025	Do not add
QAPage	Retained	Available if ever useful
VideoObject	Active	Add when the YouTube demo exists
AggregateRating	Active	Forbidden here - no real reviews exist

8. Constraints that shape everything

1. **Pre-launch, waitlist-only, nothing downloadable.** This is the dominant constraint. It makes purchase-intent and "recommend a tool" queries unwinnable, and it blocks the highest-authority directories (§9).
 2. **Effectively no backlinks.** The real reason for position 15.8.
 3. **Almost no entity presence.** AI engines weight presence on YouTube, Reddit, LinkedIn and Wikipedia. We have essentially none - and **three of those four are free and available pre-launch.**
 4. **The brand name carries a search tax.** Bing responds to "servey" with *"Do you mean **survey**?"*. Google brand search works (131 impressions), so it is not fatal, but every search box and answer engine will fight the name. Worth knowing before blaming rankings for it.
 5. **hero-devices.png** **exposes personal filenames** and is live on the homepage. The owner's decision is to leave it. Not an SEO issue; recorded so nobody "fixes" it by accident.
-

9. Channels - available vs blocked pre-launch

Verified against each platform's own rules.

Channel	Status	Note
AlternativeTo	 Blocked	<i>"If the app is Coming Soon (announced only), in closed/private beta, invite-only, or in Early Access phase, we do not accept it."</i> Account is created and past the 7-day age requirement, so submission can go in on launch day.
Show HN	 Blocked	<i>"If your work isn't ready for users to try out, please don't do a Show HN."</i>
App Store listing	 At launch	#1 most-cited source in this category (77 citations)
YouTube demo	 Available	#2 most-cited (74 citations). The highest-leverage action available today, and still not done.

See `docs/DEMO-VIDEO.md` for what the existing footage can and cannot be used for, and the shot list for recording it properly. | **Reddit** (r/macapps, r/homelab) |  Available | **#3 most-cited** (65 citations). Genuine participation, never a launch post | | **BetaList / Peerlist / Uned** |  Available | Built specifically for unlaunched products | | **HN as a normal submission** |  Available | Show HN is barred; submitting a post as a regular link is fine | | **Email the waitlist** |  Available | Still never used | | **Product Hunt** | Done | Badge in footer. Never solicit upvotes |

10. Measurement

Tool	What it answers	How
<code>npm run seo:audit</code>	Do the on-page invariants still hold?	Titles ≤60, descriptions 110–160, canonical self-reference, og:url agreement, one h1, no skipped headings, valid JSON-LD with resolving <code>@id</code> s, <code>BreadcrumbList count == visible trail</code> , all sitemap URLs 200. Exits non-zero - can gate a deploy.
Google Search Console	Impressions, position, query clusters	The cluster table in §2 is the useful view, not the top-line number
Bing / ChatGPT	Are we retrievable by answer engines?	Phrase search on Bing; test comparison queries, not "recommend me a tool"
PostHog	What happens after the click	<code>\$pageview</code> , <code>waitlist_submitted</code> → <code>waitlist_signup</code> (with <code>duplicate</code>), <code>waitlist_failed</code> , all carrying <code>source</code>
<code>npm run indexnow</code>	Push new URLs to Bing	Run after every content deploy

Reading GSC correctly

- **"Discovered – currently not indexed"** → Google knows the URL, hasn't crawled it. Requesting indexing helps.
- **"Crawled – currently not indexed"** → Google crawled it and *chose* not to index. Requesting indexing will **not** help. That is an authority verdict, and the correct response is links and entity presence, not more pages.

11. What to do next, in priority order

1. **Record the YouTube demo.** Highest-leverage action available pre-launch, feeds YouTube + Reddit + the launch-day AlternativeTo submission, and is the only top-3 cited source reachable before shipping.
2. **Reddit participation** in r/macapps and r/homelab - genuine, not promotional.
3. **Email the waitlist.** Costs nothing, never used.
4. **BetaList / Peerlist / Unneed** submissions.
5. **Add `VideoObject` schema** once the demo exists.
6. Re-test AI citation on **comparison** queries in ~2 weeks.
7. At launch: App Store listing, AlternativeTo submission, Show HN.

Explicitly not on this list: more landing pages, more metadata rewrites, more schema. On-site SEO has reached sharply diminishing returns. What is missing is that **nobody has ever seen Servey move.**

12. Change log

Date	Change
2026-09-01	Audited against the eight durable SEO strategies; two real defects found and fixed, one standing diagnosis corrected. (1) Per-query position enabled for the first time - the site is bimodal, comparison cluster at 2.3-7.9 (page one) and core product cluster at 40-44.6 (page four). The 16.5 average we have reasoned from since 14 Aug is a midpoint between two unrelated populations. §1 core diagnosis rewritten; the §4 ban on CTR/metadata work is partially lifted, because it was justified with that bad average. (2) Internal links: 2 true orphans. Simulating <code>relatedPosts()</code> rather than counting raw keyword overlap showed <code>screens-vs-jump-desktop</code> - our best non-brand page, position 7.0 - receiving zero inbound internal links, because a top-3-by-overlap rule lets hubs crowd out mid-tier pages. Raising the count to 4 or 5 does not fix it and concentrates hubs further (max inbound 7 -> 10), so fixed at keyword level: <code>Screens vs Jump Desktop</code> added to <code>screens-5-alternatives</code> (30 Screens / 12 Jump Desktop mentions) and <code>connect to Mac behind CGNAT</code> to <code>headless-mac-mini-setup</code> (has a full section on CGNAT). All 23 posts now have >= 1 inbound; min 1, max 7. (3) Index bloat: none. 28 indexed, "Crawled - not indexed" 0. (4) Keyword cannibalization: checked and not present - <code>keywords</code> drives <code>relatedPosts()</code> and inert <code><meta keywords></code>, so 9 posts sharing <code>best remote desktop for Mac</code> is a hub, not competition. Real cannibalization would be two pages targeting one query in title/H1, and the <code>/control-mac-from-iphone</code> vs <code>blog/control-your-mac-from-iphone-ipad</code> pair is deliberately split commercial/informational.
2026-09-01	Survey framing regression on the 31 Aug posts, caught by the owner. The rule in §3 (first mention <=15% depth, density 4.0-6.0/1k) was written on 26 Aug in response to this exact complaint, and the two new posts broke it: <code>control-iphone-from-mac</code> 77% depth / 3.2 per 1k, <code>headless-mac-mini-setup</code> 76% / 1.8 - less than half the density floor. The specific failure the owner quoted was the phrase pattern "a VPN or an app that handles the networking, such as Survey": Survey demoted to a trailing example of a category, listed after the alternative, in a subordinate clause. Cause worth remembering: the §2 finding that day was that the old core page failed <i>because it pitched instead of answering</i>, so I over-applied "lead with the free built-in answer" and let it suppress Survey everywhere, including in FAQ answers where it does not apply at all. Leading with the honest free option is a rule about the <i>opening frame of an informational section</i>, not a licence to bury the product. Fixed to 6% / 5.2, 4% / 4.6, 3% / 5.2, all passing. Every honesty caveat verified intact afterwards (Survey does not control iPhones; Tailscale genuinely good; Apple-only; cannot be self-hosted; Screen Sharing free and already installed). Check depth and density before committing a post, not after.
2026-08-31	Link-graph fix on arrival. <code>headless-mac-mini-setup</code> shipped with 1 inbound post, the same failure mode logged on 23 Aug. Added <code>headless Mac Mini remote control</code> as a 6th keyword to <code>real-terminal-on-your-mac-from-iphone</code> (4 headless mentions, genuinely about driving one from a shell) and <code>Mac Mini home server</code> to <code>who-is-survey-for-developers-home-labs</code> (4 Mac mini mentions). Inbound 1 -> 3. <code>control-iphone-from-mac</code> arrived with 3 and needed nothing. Reminder for future posts: check inbound overlap <i>before</i> committing, not after - keyword overlap is the entire internal link graph (§6).
2026-08-31	GSC query count went 39 -> 162, which exposed the real problem. Pulled the full query list rather than the top 10 and clustered it (§2). The <code>control my Mac from my iPhone</code> family - Survey's exact job - was ~120 impressions across 55 phrasings, converting zero, the largest cluster on the site. Cause was legible on the page: <code>control-your-mac-from-iphone-ipad</code> was a pitch, not a guide, and never named System Settings, Screen Sharing or VNC while the queries were overwhelmingly <code>how to ...</code>. Rebuilt it as a genuine how-to that leads with Apple's free built-in route, then places Survey against the away-from-home case it does not solve - the same shape that made <code>does-mac-screen-sharing-work-over-the-internet</code> convert. Added a 4-row approaches table, 7 FAQs replacing 3 pitch-only ones, <code>metaTitle</code>,

Date	Change
	and 7 keywords covering the actual phrasings. This page is the <i>informational</i> half of the pair; <code>/control-mac-from-iphone</code> stays the commercial half, so the rewrite reduces overlap rather than creating it.
2026-08-31	Two posts, 32 -> 34 URLs, both chosen from zero-click clusters rather than by adding more competitor comparisons. <code>control-iphone-from-mac</code> serves inverse intent - ~46 impressions across 13 phrasings of people wanting their <i>iPhone</i> from their <i>Mac</i> , which we had no page for and which was diluting the core page. Answered honestly with iPhone Mirroring (requirements verified against Apple support 120421) including the limitation that decides the question: the phone must be locked and <i>near</i> the Mac, so it is not remote access, and iOS permits no third-party equivalent. States plainly that Survey does not control iPhones before pivoting to the direction that does work. <code>headless-mac-mini-setup</code> serves ~32 impressions of <i>setup</i> intent (<code>headless mac mini setup</code> , <code>macos headless mode</code>) that the existing access-oriented post does not answer; real commands, and the FileVault-vs-auto-login trade-off stated as a trade-off rather than a recommendation. Link-worthy for the r/homelab outreach in BACKLINKS §2.
2026-08-31	Survey row highlighting. The green-border treatment added 27 Aug only detected a Survey <i>column</i> . Tables comparing <i>approaches</i> (rather than products) put Survey in a row, so the rebuilt core post got no highlight. Extended <code>app/blog/[slug]/page.tsx</code> to detect either orientation.
2026-08-31	Not written: a <code>screens 5 cost</code> post (~17 impressions, pure commercial intent, uncovered). Edovia does not publish pricing on the product page, and shipping figures we cannot verify - or vague copy that ranks for nothing - both fail. Revisit only with a verifiable source.
2026-08-26	Survey raised in the two Apple-tools posts. First-mention depth was 63% (<code>what-replaced-back-to-my-mac</code>) and 56% (<code>does-mac-screen-sharing-work-over-the-internet</code>), against 3% on <code>jump-desktop-vs-rustdesk</code> and 12% on <code>rustdesk-vs-anydesk</code> . The comparison posts score well because they lead with a table carrying a Survey column; the informational posts had tables with no Survey in them. Fixed structurally rather than by adding copy: the Back to My Mac jobs table gained a Survey column plus a terminal row (and an honest No on mounting the disk in the Finder), and the Screen Sharing approaches table - whose shape is approach x cost, so a Survey column would be a category error - instead names Screens, Jump Desktop, RustDesk and Survey inside the row they all belong to. Both posts now open the short-answer section by naming all four routes with a <code>which is the app we make</code> disclosure. 63% -> 14% and 56% -> 14% ; density 5.9 and 6.6 per 1k. The disclosure and the honest No are load-bearing here, not decoration: these are the pages the AppleInsider / MacRumors / 9to5Mac broken-link outreach points at.
2026-08-26	Two posts, 30 -> 32 URLs. <code>what-replaced-back-to-my-mac</code> exists to serve the broken-link campaign: AppleInsider, MacRumors and 9to5Mac all cite <code>support.apple.com/en-us/HT208922</code> , which still 404s (re-verified today, along with the other three dead Apple URLs). Our only replacement resource was <code>does-mac-screen-sharing-work-over-the-internet</code> , which covers Back to My Mac in a single section - Back to My Mac appeared 4 times sitewide, all inside that one post. An exact-match page makes the swap an easier yes for an editor. <code>rustdesk-vs-anydesk</code> feeds the only cluster earning non-brand clicks; RustDesk positions itself against AnyDesk explicitly, and the open-source axis already converts via <code>jump-desktop-vs-rustdesk</code> . Link graph engineered rather than incidental: <code>Back to My Mac replacement</code> added as a 6th keyword to the Screen Sharing post (it genuinely has a section and an FAQ on it), giving inbound counts of 6 and 3 on arrival against the 1 the previous post shipped with. Includes the AnyDesk 2024 production-systems incident, stated with its disclosure and cert rotation, because the comparison axis is literally auditable-vs-trust-the-vendor and omitting it would be dishonest. seo:audit 32/32.
2026-08-23	<code>does-mac-screen-sharing-work-over-the-internet</code> - first post in the Apple-tools cluster. Apple Remote Desktop appeared 0 times anywhere on the site; Screen Sharing and VNC appeared only as passing mentions inside comparison posts, with one keyword entry sitewide (<code>VNC alternative Mac</code>).

Date	Change
	No page targeted Apple's built-in tooling. Targets the Screen Sharing variant of the query (higher volume than the ARD variant Macky targets) and covers ARD and Remote Management inside it. Link graph: shipped with 1 inbound, so <code>macOS Screen Sharing</code> was added as a 6th keyword to the three posts that genuinely discuss it (<code>chrome-remote-desktop-vs-jump-desktop</code> 7 mentions, <code>best-remote-desktop-for-mac</code> 6, <code>screens-5-alternatives</code> 5) - inbound 1 -> 2, outbound now 3 topically-matched posts. 29 -> 30 URLs.
2026-08-23	Comparison posts re-pitched. The 6 comparison/roundup posts read as neutral review-site content: <code>jump-desktop-vs-teamviewer</code> had 2 Servey mentions in 1,526 words (1.3/1k), <code>anydesk-vs-teamviewer</code> 2 in 1,298. Their <code>Where Servey fits</code> sections actively deflected ("it does not compete with either of these", "pick one of the two above") and all 5 <code>Bottom line</code> sections had no Servey at all. Added a Servey column to all 5 head-to-head tables (each already ended with a `Real terminal
2026-08-19	Internal-linking fix: the 7 use-case pages were orphaned - they cross-linked to each other but nothing linked in, so the homepage emitted only 4 internal links (<code>/</code> , <code>/blog</code> , <code>/privacy</code> , <code>/terms</code>). GSC counted 33 internal links sitewide with <code>/privacy</code> , <code>/terms</code> , <code>/blog</code> as top targets, i.e. the footer was the only link graph Google saw. Added a footer <code>Use cases</code> nav (homepage 4 -> 11 links) and a <code>Servey for this</code> block on posts inverting <code>relatedSlug</code> .
2026-08-19	3 comparison posts (<code>splashtop-vs-jump-desktop</code> , <code>anydesk-vs-teamviewer</code> , <code>chrome-remote-desktop-vs-jump-desktop</code>), keyword-wired to <code>screens-vs-jump-desktop</code> . 26 -> 29 URLs.
2026-08-19	Em dashes replaced with <code>-</code> across code comments and docs (owner preference). None were in rendered copy. <code>source-material/</code> left untouched as owner input.
2026-08-16	3 comparison posts (<code>screens-5-alternatives</code> , <code>jump-desktop-vs-teamviewer</code> , <code>jump-desktop-vs-rustdesk</code>); <code>/control-mac-from-iphone</code> pillar; headless + local-AI pages retargeted; Speculation Rules; image <code>sizes</code> fix; <code>ImageObject</code> JSON-LD; <code>og:image</code> dimensions on all pages; <code>/blog</code> twitter card fix; <code>Article</code> publisher by <code>@id</code> ; <code>scripts/seo-audit.mjs</code> . 22 -> 26 URLs. Sitemap re-read by Google: 26 discovered.
2026-08-16	PostHog analytics (lazy-loaded, first-party proxied, DNT honoured, no person profiles).
2026-08-14	6 use-case landing pages; Termius comparison; <code>table</code> block type; breadcrumbs; RSS feed; 404 page; light-mode contrast and reduced-motion fixes. 15 -> 22 URLs.
Earlier	JSON-LD <code>@graph</code> ; IndexNow; AI-crawler allowlist; <code>llms.txt</code> ; comparison and roundup posts.