

Servey - SEO Context

The SEO playbook for `servey.in`: what we do, what we deliberately do not do, and the evidence behind each decision.

Companion to `CONTEXT.md`, which covers the product and codebase. This file covers search and answer-engine strategy only.

- **Last reviewed:** 26 August 2026
- **Source of truth:** `docs/SEO-CONTEXT.md` . `docs/seo-context.pdf` is generated from it by `npm run context:pdf` - never edit the PDF.

Why this file exists. Nearly every conclusion below was expensive to reach and easy to forget. Several were reached *twice* because the reasoning was not written down. If you are about to "improve" something here, read the reason first - most of these are deliberate trade-offs, not oversights.

1. Baseline - where we actually are

Google Search Console, 3-month window to **25 August 2026**:

Metric	Value	11 days earlier
Clicks	38	16
Impressions	1,460	626
CTR	2.6%	2.6%
Average position	16.4	15.8

Impressions and clicks both roughly 2.3x in eleven days. This is by some way the largest movement the site has recorded, and it is the first reading where the trend is unambiguous rather than noise. CTR held exactly flat at 2.6% while volume more than doubled, which is the healthy shape: the new impressions are converting at the same rate as the old ones, so we are reaching more of the same kind of query rather than drifting into irrelevant ones.

Position drifting 15.8 -> 16.4 alongside that growth is **not** a regression. New pages enter the index low and pull the average down; a rising page count with rising clicks and a falling average position is what growth looks like from here. Judge this site on clicks and impressions, not on average position, until the page count stops changing.

Brand vs non-brand: `survey` alone is 9 clicks / 215 impressions - **24% of clicks from 15% of impressions.** This was 44% of clicks eleven days ago. Brand volume grew (7 -> 9 clicks) while its *share* nearly halved, which means the growth is coming from non-brand queries. That is the number this whole strategy was aimed at, and it moved.

Top non-brand queries: `screens vs jump desktop` (3 clicks / 22 impressions), `jump desktop vs screens 5` (1 / 19), `jump desktop vs rustdesk` (1 / 16). The competitor-vs-competitor cluster remains the only non-brand cluster earning clicks, and it is now earning more of them. Note `jump desktop vs rustdesk` appearing here for the first time - a post shipped 19 Aug now converts.

Index coverage (verified 26 Aug 2026):

- **Google** - sitemap submitted, **32 URLs**, last read 16 Aug 2026.
- **Bing** - indexed. Homepage plus at least four blog URLs return on a phrase search, and Bing's own AI summary describes Survey accurately from our pages. This matters because **ChatGPT Search reads Bing's index**, so the pipe to the largest answer engine is connected.

The core diagnosis

Position 16.4 is page two, which takes roughly **1% of clicks** versus ~25% on page one. CTR of 2.6% is *above* par for that position, which means **titles and descriptions are working**. This is a **ranking problem, not a copywriting problem**, and the bottleneck is authority - we have effectively no backlinks.

The 2.3x growth to 25 Aug does not change this diagnosis, it confirms it. We are winning more impressions because we have more relevant pages, and converting them at a steady rate. What we are still not doing is moving off page two, and no amount of additional content does that on its own.

Do not respond to flat numbers by rewriting metadata again. That lever is already pulled. The remaining levers are links, entity presence, and shipping.

2. What the query data proves

All 39 queries, clustered:

Cluster	Impressions	Clicks	Verdict
Competitor vs competitor (screens vs jump desktop + ~11 variants)	~58	2	✅ The only non-brand cluster earning clicks
AI on Mac (run local ai on mac , mac ai agent)	~15	0	Largest non-competitor cluster
Control Mac from iPhone (10 phrasings)	~12	0	Core intent - was fragmented
Headless Mac mini (7 phrasings)	~9	0	Was fragmented
airplayuiagent	4	0	Accidental, off-intent (a macOS process name)

Three conclusions that drive every content decision:

1. **Comparison content is the only proven format** - specifically *competitor-vs-competitor*, where Servey is **not** the subject. We rank by genuinely helping someone choose between two other products.
 2. **Fragmentation is the failure mode.** Ten phrasings of one intent, each with 1-2 impressions and zero clicks, is what Google does when it reads a site as marginally relevant to all of a group and authoritative on none. The fix is **one strong pillar**, not more pages.
 3. **Pure how-to content is not ranking.** It earns impressions and no clicks.
-

3. What we DO

Content

- **Write competitor-vs-competitor comparisons.** The proven format. Currently: `screens-vs-jump-desktop`, `jump-desktop-vs-teamviewer`, `jump-desktop-vs-rustdesk`, `rustdesk-vs-anydesk`, `screens-5-alternatives`, `screens-jump-desktop-alternative-mac`, `termius-alternative-mac-terminal`, `best-remote-desktop-for-mac`.
- **Lead comparison posts with a `table` block.** Tables are the structure answer engines extract most reliably.
- **State where every competitor is genuinely better.** This is the single reason these pages get read and cited. It is a strategy, not politeness.
- **Keep Servey out of the hero slot** in competitor-vs-competitor posts: one late mention, labelled pre-launch, ending by telling the reader to buy an established product if they need one today.
- **One commercial pillar per intent**, with informational blog posts as spokes pointing at it (`/control-mac-from-iphone` ← `control-your-mac-from-iphone-ipad`).

On-page

- `metaTitle` ≤51 chars (the `%s - Servey` template adds 9; rendered ≤60).
- Descriptions **110–160 chars**.
- Explicit `alternates.canonical` and explicit `openGraph.url` on every page.
- One `<h1>`, no skipped heading levels.
- Shared `ogImage` object (never the bare string, which drops width/height/alt).

Technical

- **Static generation everywhere.** Every route is Static/SSG; guard this.
- **Speculation Rules** - same-origin prefetch/prerender at `moderate` eagerness, excluding `/api/` and `/ingest/`.
- **Structured data:** one `@graph` (Organization → WebSite → SoftwareApplication with PreOrder Offers → FAQPage → ImageObject), plus per-page `Article` / `WebPage` / `BreadcrumbList` / `FAQPage`. Cross-reference by `@id`; never re-declare a node inline.
- **sizes on every responsive image**, matched to its real rendered width.
- **IndexNow** after every content deploy (`npm run indexnow`) - fastest route into Bing, and therefore into ChatGPT Search.
- `npm run seo:audit` before/after deploys. Currently **32/32 pass**.

Answer engines (AEO / GEO)

- **Self-contained passages of ~130–170 words** that still make sense quoted in isolation.
- **Question-shaped H2s** mirroring real queries.
- **Direct answer in the first sentence** after each heading, then support.
- **robots.ts explicitly allows AI crawlers by name** - GPTBot, OAI-SearchBot, ChatGPT-User, ClaudeBot, Claude-Web, Claude-SearchBot, PerplexityBot, Perplexity-User, Google-Extended, Applebot-Extended, Bingbot, Amazonbot, CCBot.

4. What we deliberately DO NOT do

This section is as important as the last one. Each of these is a considered decision with a reason.

Never - integrity

We don't	Why
Invent ratings, testimonials, review counts, user numbers, or <code>AggregateRating</code>	Survey is pre-launch with zero public users. <code>AggregateRating</code> without real reviews is a Google structured-data violation, and the rest is simply false.
Fabricate product screenshots	Every image is a real capture or an honestly-styled placeholder.
Publish a comparison that hides where a competitor wins	It is the entire credibility asset. A post that reads like an advert has failed at its job.
Imply Termius shipped malware	They were the victim of the ZuRu impersonation campaign. Saying otherwise is false and defamatory.
Solicit upvotes anywhere	Against Product Hunt / AlternativeTo rules; triggers rank penalties.
Add our own links to Wikipedia (dead-link building)	Verified 24 Aug: Wikipedia renders external links as <code>rel="nofollow"</code> sitewide, so they pass no ranking signal. Adding our own also violates WP:COI/WP:REFSPAM, where the failure mode is <code>servey.in</code> on the public Wikipedia spam blacklist. Wikipedia is downstream of independent press, not a substitute. Broken-link building off Wikipedia is fine and is our active tactic - see <code>docs/BACKLINKS.md</code> .
Buy links, use PBNs, or exchange links	Cheap short-term, fatal long-term.
Work a bulk "N places to get backlinks" list top to bottom	Triaged one of 75 on 26 Aug: 9 were usable now . The rest were editorial placements you must earn (Forbes, TechCrunch), revenue-gated directories we cannot enter pre-launch, review sites needing real customers, open-source-only listings, or content farms. Submitting to dozens of unrelated low-quality directories in a burst is the classic footprint of a link scheme, and on a domain with 2 referring domains it is the <i>only</i> pattern in our profile. Pick for relevance, submit steadily, and never post the same text everywhere without a canonical. See <code>docs/BACKLINKS.md</code> §6.
Keyword-stuff, cloak, or hide text	Trivially detected; the fragmentation fix is pillar pages, not stuffing.
Publish thin AI-generated filler to inflate page count	Google's QRG §4.6.5–6 targets exactly this. Page count is not the goal.

Not now - strategy

We don't	Why
Chase "recommend me a tool" AI queries	Verified 16 Aug: ChatGPT answered "tool to screen share to my Mac from my phone" with Chrome Remote Desktop, AnyDesk, TeamViewer. Correctly - Servey has nothing to download, so recommending it would be a worse answer. Structurally unwinnable until launch. Do not treat it as an SEO failure.
Chase "termius alternative" broadly	It would pull multi-host-SSH searchers who bounce, polluting waitlist quality pre-launch. The honest overlap is narrow: "run a terminal on your own Mac without SSH setup".
Compete on Termius's page axis (SSH client x platform)	They hold #1 there. We copied their <i>template technique</i> onto a different axis - task x device - where they do not rank at all.
Add more pages targeting "control Mac from iPhone"	That recreates the exact fragmentation <code>/control-mac-from-iphone</code> was built to fix. One pillar, many spokes.
Rewrite metadata again to chase CTR	2.6% is already above par for position 16.4, and held flat while impressions grew 2.3x. The lever is pulled.
Build doorway pages per query variant	The 10 phrasings get absorbed by one strong page, not 10 weak ones.
Do multilingual / translated SEO	Considered 24 Aug and deliberately held. 19 posts x 7 locales is 133 URLs on a domain with 2 referring domains already averaging position 16.3. Translations multiply pages, not authority, and scaled machine translation without human review is named in Google's spam policies. The tactic works for language-independent utility queries ("instagram photo downloads"); ours are niche Mac comparison queries whose product names, docs and buyers are all English. Revisit once referring domains are in double digits, then as a one-language, 4-post pilot chosen from GSC Countries data - never a 7-locale bulk translation.
Chase <code>airplayuiagent</code> traffic	Accidental, off-intent, and converts nobody.

Myths we do not act on

Belief	Reality
<code>llms.txt</code> improves AI citation	It does not. No major engine consumes it; this is evidence-based, not opinion. We keep ours because it costs nothing and is a tidy human-readable map - but it is not a citation lever and must not be cited as one.
FAQ schema wins SERP real estate	Retired 7 May 2026. See §7.
More pages = more traffic	39 queries across 26 pages says relevance is fine; authority is the constraint.
Content needs chunking / AI-specific keyword rewriting	Engines handle synonyms; this is wasted effort.

5. Cluster → page ownership map

Each intent has exactly one owner. Adding a second competitor for any row is how fragmentation comes back.

Intent cluster	Owning page	Supporting spokes
Control Mac from iPhone	/control-mac-from-iphone	blog/control-your-mac-from-iphone-ipad
Terminal from iPhone	/terminal-on-iphone	blog/real-terminal-on-your-mac-from-iphone
Terminal from iPad	/terminal-on-ipad	-
Headless Mac mini / macOS headless mode	/headless-mac-mini	blog/control-a-headless-mac-mini-remotely
Run local AI on Mac / mac ai agent	/remote-mac-for-ai-agents	blog/run-ai-agents-locally-on-your-mac , blog/run-ai-agents-on-your-mac-remotely , blog/stay-in-control-of-ai-agents-from-anywhere
Developer remote Mac access	/mac-for-developers	blog/who-is-servey-for-developers-home-labs
Mac home lab	/mac-home-lab	-
Cellular / strict networks	blog/access-your-mac-remotely-over-cellular	-
Screens vs Jump Desktop	blog/screens-vs-jump-desktop	blog/screens-jump-desktop-alternative-mac
Back to My Mac replacement	blog/what-replaced-back-to-my-mac	blog/does-mac-screen-sharing-work-over-the-internet
RustDesk vs AnyDesk	blog/rustdesk-vs-anydesk	-
Screens 5 alternatives	blog/screens-5-alternatives	-
Jump Desktop vs TeamViewer	blog/jump-desktop-vs-teamviewer	-
Jump Desktop vs RustDesk	blog/jump-desktop-vs-rustdesk	-
Termius / SSH overlap	blog/termius-alternative-mac-terminal	-
Category roundup	blog/best-remote-desktop-for-mac	-

6. Internal linking - a quirk that matters

`p` blocks render as plain text. `app/blog/[slug]/page.tsx` parses no markdown and no HTML, so **you cannot put an anchor in body copy**. The only real internal-link mechanism between posts is **exact-string keyword overlap** feeding `relatedPosts()`.

Engineering the `keywords` array *is* engineering the link graph. When adding a post, deliberately share keywords with the posts you want linking to it. The three newest comparison posts were given keyword overlap with `screens-vs-jump-desktop` specifically because it is the one page earning clicks, so its authority flows to them.

Use-case pages *do* have real anchors: the cross-link mesh is derived automatically from the `useCases` array order, so adding a page wires it in.

7. Structured-data deprecation register

Type	Status	Action
FAQPage	Rich results retired 7 May 2026. SC report + Rich Results Test support removed June 2026; API data August 2026.	Keep the markup. It remains a valid Schema.org type, Google confirms unused structured data does not affect Search, and answer engines still read it. Only the <i>expectation</i> changed.
HowTo	Removed June 2025	Do not add
SpecialAnnouncement , ClaimReview , VehicleListing , EstimatedSalary , LearningVideo , Course carousel, Book Actions	Removed June 2025	Do not add
QAPage	Retained	Available if ever useful
VideoObject	Active	Add when the YouTube demo exists
AggregateRating	Active	Forbidden here - no real reviews exist

8. Constraints that shape everything

1. **Pre-launch, waitlist-only, nothing downloadable.** This is the dominant constraint. It makes purchase-intent and "recommend a tool" queries unwinnable, and it blocks the highest-authority directories (§9).
 2. **Effectively no backlinks.** The real reason for position 15.8.
 3. **Almost no entity presence.** AI engines weight presence on YouTube, Reddit, LinkedIn and Wikipedia. We have essentially none - and **three of those four are free and available pre-launch.**
 4. **The brand name carries a search tax.** Bing responds to "servey" with *"Do you mean **survey**?"*. Google brand search works (131 impressions), so it is not fatal, but every search box and answer engine will fight the name. Worth knowing before blaming rankings for it.
 5. **hero-devices.png** **exposes personal filenames** and is live on the homepage. The owner's decision is to leave it. Not an SEO issue; recorded so nobody "fixes" it by accident.
-

9. Channels - available vs blocked pre-launch

Verified against each platform's own rules.

Channel	Status	Note
AlternativeTo	 Blocked	<i>"If the app is Coming Soon (announced only), in closed/private beta, invite-only, or in Early Access phase, we do not accept it."</i> Account is created and past the 7-day age requirement, so submission can go in on launch day.
Show HN	 Blocked	<i>"If your work isn't ready for users to try out, please don't do a Show HN."</i>
App Store listing	 At launch	#1 most-cited source in this category (77 citations)
YouTube demo	 Available	#2 most-cited (74 citations). The highest-leverage action available today, and still not done.
Reddit (r/macapps, r/homelab)	 Available	#3 most-cited (65 citations). Genuine participation, never a launch post
BetaList / Peerlist / Uneed	 Available	Built specifically for unlaunched products
HN as a normal submission	 Available	Show HN is barred; submitting a post as a regular link is fine
Email the waitlist	 Available	Still never used
Product Hunt	Done	Badge in footer. Never solicit upvotes

10. Measurement

Tool	What it answers	How
<code>npm run seo:audit</code>	Do the on-page invariants still hold?	Titles ≤60, descriptions 110–160, canonical self-reference, og:url agreement, one h1, no skipped headings, valid JSON-LD with resolving <code>@id</code> s, <code>BreadcrumbList count == visible trail</code> , all sitemap URLs 200. Exits non-zero - can gate a deploy.
Google Search Console	Impressions, position, query clusters	The cluster table in §2 is the useful view, not the top-line number
Bing / ChatGPT	Are we retrievable by answer engines?	Phrase search on Bing; test comparison queries, not "recommend me a tool"
PostHog	What happens after the click	<code>\$pageview</code> , <code>waitlist_submitted</code> → <code>waitlist_signup</code> (with <code>duplicate</code>), <code>waitlist_failed</code> , all carrying <code>source</code>
<code>npm run indexnow</code>	Push new URLs to Bing	Run after every content deploy

Reading GSC correctly

- **"Discovered – currently not indexed"** → Google knows the URL, hasn't crawled it. Requesting indexing helps.
- **"Crawled – currently not indexed"** → Google crawled it and *chose* not to index. Requesting indexing will **not** help. That is an authority verdict, and the correct response is links and entity presence, not more pages.

11. What to do next, in priority order

1. **Record the YouTube demo.** Highest-leverage action available pre-launch, feeds YouTube + Reddit + the launch-day AlternativeTo submission, and is the only top-3 cited source reachable before shipping.
2. **Reddit participation** in r/macapps and r/homelab - genuine, not promotional.
3. **Email the waitlist.** Costs nothing, never used.
4. **BetaList / Peerlist / Unneed** submissions.
5. **Add VideoObject schema** once the demo exists.
6. Re-test AI citation on **comparison** queries in ~2 weeks.
7. At launch: App Store listing, AlternativeTo submission, Show HN.

Explicitly not on this list: more landing pages, more metadata rewrites, more schema. On-site SEO has reached sharply diminishing returns. What is missing is that **nobody has ever seen Servey move.**

12. Change log

Date	Change
2026-08-26	Two posts, 30 -> 32 URLs. <code>what-replaced-back-to-my-mac</code> exists to serve the broken-link campaign: AppleInsider, MacRumors and 9to5Mac all cite <code>support.apple.com/en-us/HT208922</code> , which still 404s (re-verified today, along with the other three dead Apple URLs). Our only replacement resource was <code>does-mac-screen-sharing-work-over-the-internet</code> , which covers Back to My Mac in a single section - Back to My Mac appeared 4 times sitewide, all inside that one post. An exact-match page makes the swap an easier yes for an editor. <code>rustdesk-vs-anydesk</code> feeds the only cluster earning non-brand clicks; RustDesk positions itself against AnyDesk explicitly, and the open-source axis already converts via <code>jump-desktop-vs-rustdesk</code> . Link graph engineered rather than incidental: <code>Back to My Mac replacement</code> added as a 6th keyword to the Screen Sharing post (it genuinely has a section and an FAQ on it), giving inbound counts of 6 and 3 on arrival against the 1 the previous post shipped with. Includes the AnyDesk 2024 production-systems incident, stated with its disclosure and cert rotation, because the comparison axis is literally auditable-vs-trust-the-vendor and omitting it would be dishonest. seo:audit 32/32.
2026-08-23	<code>does-mac-screen-sharing-work-over-the-internet</code> - first post in the Apple-tools cluster. Apple Remote Desktop appeared 0 times anywhere on the site; Screen Sharing and VNC appeared only as passing mentions inside comparison posts, with one keyword entry sitewide (<code>VNC alternative Mac</code>). No page targeted Apple's built-in tooling. Targets the Screen Sharing variant of the query (higher volume than the ARD variant Macky targets) and covers ARD and Remote Management inside it. Link graph: shipped with 1 inbound, so <code>macOS Screen Sharing</code> was added as a 6th keyword to the three posts that genuinely discuss it (<code>chrome-remote-desktop-vs-jump-desktop</code> 7 mentions, <code>best-remote-desktop-for-mac</code> 6, <code>screens-5-alternatives</code> 5) - inbound 1 -> 2, outbound now 3 topically-matched posts. 29 -> 30 URLs.
2026-08-23	Comparison posts re-pitched. The 6 comparison/roundup posts read as neutral review-site content: <code>jump-desktop-vs-teamviewer</code> had 2 Servey mentions in 1,526 words (1.3/1k), <code>anydesk-vs-teamviewer</code> 2 in 1,298. Their <code>Where Servey fits</code> sections actively deflected ("it does not compete with either of these", "pick one of the two above") and all 5 <code>Bottom line</code> sections had no Servey at all. Added a Servey column to all 5 head-to-head tables (each already ended with a `Real terminal
2026-08-19	Internal-linking fix: the 7 use-case pages were orphaned - they cross-linked to each other but nothing linked in, so the homepage emitted only 4 internal links (<code>/</code> , <code>/blog</code> , <code>/privacy</code> , <code>/terms</code>). GSC counted 33 internal links sitewide with <code>/privacy</code> , <code>/terms</code> , <code>/blog</code> as top targets, i.e. the footer was the only link graph Google saw. Added a footer <code>Use cases</code> nav (homepage 4 -> 11 links) and a <code>Servey for this</code> block on posts inverting <code>relatedSlug</code> .
2026-08-19	3 comparison posts (<code>splashtop-vs-jump-desktop</code> , <code>anydesk-vs-teamviewer</code> , <code>chrome-remote-desktop-vs-jump-desktop</code>), keyword-wired to <code>screens-vs-jump-desktop</code> . 26 -> 29 URLs.
2026-08-19	Em dashes replaced with <code>-</code> across code comments and docs (owner preference). None were in rendered copy. <code>source-material/</code> left untouched as owner input.
2026-08-16	3 comparison posts (<code>screens-5-alternatives</code> , <code>jump-desktop-vs-teamviewer</code> , <code>jump-desktop-vs-rustdesk</code>); <code>/control-mac-from-iphone</code> pillar; headless + local-AI pages retargeted; Speculation Rules; image <code>sizes</code> fix; <code>ImageObject</code> JSON-LD; <code>og:image</code> dimensions on all pages; <code>/blog</code> twitter card fix; <code>Article</code> publisher by <code>@id</code> ; <code>scripts/seo-audit.mjs</code> . 22 -> 26 URLs. Sitemap re-read by Google: 26 discovered.
2026-08-16	PostHog analytics (lazy-loaded, first-party proxied, DNT honoured, no person profiles).

Date	Change
2026-08-14	6 use-case landing pages; Termius comparison; <code>table</code> block type; breadcrumbs; RSS feed; 404 page; light-mode contrast and reduced-motion fixes. 15 → 22 URLs.
Earlier	JSON-LD <code>@graph</code> ; IndexNow; AI-crawler allowlist; <code>llms.txt</code> ; comparison and roundup posts.