

Servey - SEO Context

The SEO playbook for servey.in: what we do, what we deliberately do not do, and the evidence behind each decision.

Companion to `CONTEXT.md`, which covers the product and codebase. This file covers search and answer-engine strategy only.

- **Last reviewed:** 1 September 2026
- **Source of truth:** `docs/SEO-CONTEXT.md` . `docs/seo-context.pdf` is generated from it by `npm run context:pdf` - never edit the PDF.

Why this file exists. Nearly every conclusion below was expensive to reach and easy to forget. Several were reached *twice* because the reasoning was not written down. If you are about to "improve" something here, read the reason first - most of these are deliberate trade-offs, not oversights.

1. Baseline - where we actually are

Google Search Console, 3-month window to **29 August 2026** (read 31 Aug):

Metric	Value	25 Aug	14 Aug
Clicks	54	38	16
Impressions	2,200	1,460	626
CTR	2.5%	2.6%	2.6%
Average position	16.5	16.4	15.8
Distinct queries	162	39	-

Clicks +42% and impressions +51% in four days; the query count more than quadrupled. The query count is the number that matters most here. Going from 39 to 162 distinct queries means Google has stopped treating us as a handful of pages about Jump Desktop and started matching us across the whole category.

Position holding flat at ~16.5 while impressions grow is the expected shape and is **not** a regression - see the note in §10 on reading position correctly.

Brand vs non-brand: `survey` is 9 clicks / 257 impressions - **17% of clicks from 12% of impressions.**

Brand share of clicks has gone 44% -> 24% -> 17% over three readings while brand click volume stayed flat at 9. Non-brand is now carrying essentially all growth. That was the goal of this whole strategy.

Attribution caveat: the 162-row query table accounts for only **23 of the 54 clicks**. GSC anonymises rare queries, so ~31 clicks are unattributable to a specific term. Cluster click counts below are therefore floors, not totals.

Index coverage (verified 31 Aug 2026):

- **Google** - sitemap submitted, now **40 URLs** (28 posts, 7 use cases, `/mac` , `/` , `/blog` , `/privacy` , `/terms`).
- **Bing** - the blocking status is **Discovered but not crawled** , not "not indexed". Two corrections landed here on 1 Sep, in this order:
 1. The original note said "indexed... so that pipe is connected." Never verified. Wrong.
 2. The replacement said "9 URLs indexed of 34, not one comparison post among them," taken from **Site Explorer**. **Also wrong** - Site Explorer undercounts. `screens-vs-jump-desktop` returns "**Indexed successfully**" under URL Inspection while being **entirely absent from Site Explorer**, on both the `Indexed URLs` and `All URLs` filters.
- **Site Explorer is not the index. Use URL Inspection** (left nav) for any per-URL question. Site Explorer is a lagging, partial report and must not be used to conclude a URL is missing.
- **Full per-URL audit of the blog, 1 Sep.** Every post checked through URL Inspection; **16 Request indexing submissions made** (quota went 99 -> 83, confirming all 16 landed - the daily cap is **100**, not 10):

URL	Status before	Discovered	Submitted
blog/splashtop-vs-jump-desktop	Discovered, not crawled	19 Aug	yes
blog/jump-desktop-vs-rustdesk	Discovered, not crawled	31 Aug	yes
blog/chrome-remote-desktop-vs-jump-desktop	Discovered, not crawled	19 Aug	yes
blog/jump-desktop-vs-teamviewer	Discovered, not crawled	31 Aug	yes
blog/screens-5-alternatives	Discovered, not crawled	31 Aug	yes
blog/control-iphone-from-mac	Discovered, not crawled	31 Aug	yes
blog/headless-mac-mini-setup	Discovered, not crawled	31 Aug	yes
blog/anydesk-vs-teamviewer	Discovered, not crawled	(null)	yes
blog/rustdesk-vs-anydesk	Discovered, not crawled	26 Aug	yes
blog/termius-alternative-mac-terminal	Discovered, not crawled	16 Aug	yes
blog/does-mac-screen-sharing-work-over-the-internet	Discovered, not crawled	24 Aug	yes
blog/what-replaced-back-to-my-mac	Discovered, not crawled	(null)	yes
blog/control-your-mac-from-iphone-ipad	Indexed (pre-rebuild copy)	-	yes, recrawl
blog/screens-jump-desktop-alternative-mac	Indexed	-	yes, recrawl
blog/access-your-mac-remotely-over-cellular	Indexed	-	yes, recrawl
blog/who-is-servey-for-developers-home-labs	Indexed	-	yes, recrawl
blog/screens-vs-jump-desktop	Indexed	-	no, unchanged
blog/best-remote-desktop-for-mac	Indexed	-	no, unchanged

- **So Bing's real coverage was better than Site Explorer implied but worse than Google's:** roughly 6 of 18 posts indexed, the rest sitting uncrawled. Note `control-your-mac-from-iphone-ipad` was indexed as the **old pitch version**, pre-31-Aug rebuild - a recrawl was requested for exactly that reason.
- Several rows report `Discovered on 01 Jan 2006`, which is a null placeholder, not a real date. Do not read anything into it.
- **Discovered but not crawled is the tractable bucket.** Bing knows the URL and has not fetched it yet. This is *not* the equivalent of Google's `Crawled - currently not indexed`, which is a quality verdict; it is a crawl budget queue, and **Request indexing is the correct lever for it**. Bing's own copy on the panel says so.
- **Consequence still stands, with a smaller number:** ChatGPT Search reads Bing's index, so every comparison post stuck in this bucket is invisible to it. A distribution gap, not a content gap.
- Bing's sitemap read was last crawled **19 Aug (29 URLs)**, 5 behind live. IndexNow re-push of all 34 sent and accepted 1 Sep - note this alone did **not** move these URLs out of `Discovered but not crawled`, which is why per-URL Request indexing is needed on top of it.

The core diagnosis - CORRECTED 1 Sep 2026

The site is bimodal, and the 16.5 average is a meaningless midpoint between two populations that behave completely differently. Per-query position data pulled 1 Sep (the metric was never enabled before, which is how this went unnoticed for three readings):

Query	Position	Impressions	Clicks
<code>splashtop vs jump desktop</code>	2.3	9	1
<code>jump desktop vs rustdesk</code>	3.3	25	1
<code>jump desktop vs chrome remote desktop</code>	3.3	19	1
<code>rustdesk vs jump desktop</code>	3.8	17	4
<code>jump desktop vs splashtop</code>	4.0	8	1
<code>servey (brand)</code>	5.0	257	9
<code>screens vs jump desktop</code>	7.0	27	3
<code>jump desktop vs screens 5</code>	7.9	24	1
<code>screens 5 vs jump desktop</code>	10.8	24	1
<code>remote control iphone from mac</code>	44.6	23	0
<code>remote access mac from iphone</code>	40.0	18	0

The comparison cluster is already on page one - top five for most of it. The core product cluster is on page four to five. Averaging those two gives 16.5, a number that describes nothing real and that we have been reasoning from since 14 August.

What this changes:

1. **"We are stuck on page two, it is an authority problem" was wrong as stated.** We are on page one where we have topical depth (six Jump Desktop comparisons) and on page four where we have one page. Authority is still the binding constraint on the *core* cluster, but the comparison cluster proves the site can rank when the topic is covered properly.
2. **The §4 ruling against metadata work was based on the bad average and is now partially lifted** - see that section. CTR at position 3 is not evaluated against the same benchmark as CTR at position 16.
3. **The core-cluster rebuild of 31 Aug cannot be judged on CTR.** At position 40 the page is not being seen, so a flat click count over the next fortnight means nothing. Watch **position**, not clicks, for that page.

Index coverage (verified 1 Sep 2026): 28 indexed, 3 not indexed (2 "Page with redirect", 1 "Discovered - not indexed"). **"Crawled - currently not indexed" is 0**, which is the bucket that carries a quality verdict. There is **no index bloat** and nothing has been rejected on quality grounds.

2. What the query data proves

All 162 queries, clustered (3-month window to 29 Aug 2026):

Cluster	Phrasings	Impressions	Clicks	Verdict
Competitor vs competitor (Jump Desktop is in nearly all of them)	~20	~235	14	✅ Earns every attributable non-brand click
Control Mac FROM iPhone (remote access mac from iphone , control mac with iphone , ...)	55	~120	0	❌ Our core product intent. Biggest miss on the site
Control iPhone FROM Mac (remote control iphone from mac , ...)	13	~46	0	❌ Inverse intent - we had no page for it at all
Headless Mac mini (headless mac mini setup , macos headless mode , ...)	12	~32	0	❌ How-to intent; we only had a "why" page
AI on Mac (mac ai agent , run local ai on mac)	10	~26	0	Three posts already; saturated, do not add more
Screens 5 pricing (screens 5 cost , screens 5 price)	4	~17	0	Commercial intent, uncovered. Blocked: pricing not verifiable
airplayuiagent , publicrelay , splashbuddy	5	~8	0	Accidental - macOS process names. Ignore

The finding that reset priorities on 31 Aug: the "control my Mac from my iPhone" family - Servey's exact job - produces ~120 impressions across 55 distinct phrasings and converts zero of them. It is by far the largest cluster on the site and the only one where we should be the definitive answer.

The cause was diagnosable by reading the page. `control-your-mac-from-iphone-ipad` was written as a *pitch* ("what good Mac remote control should feel like"), not as a *guide*. It never mentioned System Settings, Screen Sharing, VNC or any actual step. Google ranked it for the topic and then buried it, because the queries are overwhelmingly `how to ...` and the page answered nothing. Rebuilt 31 Aug as a real how-to that leads with Apple's free built-in route.

Four conclusions that drive every content decision:

1. **Comparison content is the only proven format** - specifically *competitor-vs-competitor*, where Servey is **not** the subject. We rank by genuinely helping someone choose between two other products.
2. **Fragmentation is the failure mode.** 55 phrasings of one intent with zero clicks is what Google does when it reads a site as marginally relevant to all of a group and authoritative on none.
3. **A page that pitches instead of answering cannot rank for a how-to query**, however well optimised. Lead with the free/built-in answer, then position Servey against the part it genuinely does not solve. This is exactly the shape that made `does-mac-screen-sharing-work-over-the-internet` convert.
4. **Check for inverse intent.** ~46 impressions were people wanting the *opposite* direction (their iPhone from their Mac). Serving that on its own page both captures the traffic and stops it diluting the core page.

3. What we DO

Content

- **Write competitor-vs-competitor comparisons.** The proven format. Currently: `screens-vs-jump-desktop`, `jump-desktop-vs-teamviewer`, `jump-desktop-vs-rustdesk`, `rustdesk-vs-anydesk`, `screens-5-alternatives`, `screens-jump-desktop-alternative-mac`, `termius-alternative-mac-terminal`, `best-remote-desktop-for-mac`.
- **Lead comparison posts with a `table` block.** Tables are the structure answer engines extract most reliably.
- **State where every competitor is genuinely better.** This is the single reason these pages get read and cited. It is a strategy, not politeness.
- **Keep Servey out of the hero slot** in competitor-vs-competitor posts: one late mention, labelled pre-launch, ending by telling the reader to buy an established product if they need one today.
- **One commercial pillar per intent**, with informational blog posts as spokes pointing at it (`/control-mac-from-iphone` ← `control-your-mac-from-iphone-ipad`).
- **Name Servey in the first ~15% of the article, and measure it with `npm run build && npm run seo:framing` (`scripts/servey-depth.mjs`),** which runs against rendered HTML because reading order is the thing being measured. **Three** metrics, not two:
 - **first-mention depth ≤15%** - where the word first appears.
 - **substance ≤35%** (*added 5 Sep*) - where the first real paragraph (≥25 words) or `<h2>` naming Servey starts, **excluding the lede**. Depth alone is satisfiable by a table cell and was satisfied on all four 5 Sep posts, which the owner still read as burying Servey: depth 4.9-9.0% while the first Servey-bearing section sat at 49-69%. **Depth says the word is early; substance says the reader is.**
 - **density 4.0-6.0 per 1k words** (posts under 700 words are allowed up to 9.0 - a short post runs over mechanically). These apply to **informational** posts. **Comparison and roundup posts are exempt from depth and substance by design** - §3 says keep Servey out of the hero slot there - so the script classifies by slug and only reports those two for them. Do not "fix" a comparison post's late Servey section. The fix when substance is deep is **a short-answer paragraph after the opening table** that answers the question and names Servey in it with an explicit "the app we make" disclosure - not more mentions. Trimming words while keeping mentions pushes density up and reads as an advert. The informational posts both drifted to 56% and 63% because they had no reason to name Servey until the dedicated section - a reader on a phone scrolls a long way before learning a product exists. The fix that works is **structural, not a pitch**: give the post's table a Servey column, or name Servey inside the row or list item it genuinely belongs to, alongside the competitors.
- **On pages used for outreach, keep the honest "No".** The Back to My Mac table says Servey does *not* mount your Mac's disk in the Finder, because it doesn't. A table that only says yes is worth nothing to a reader, and an editor deciding whether to swap a dead citation for our page will not link to an advert.

On-page

- `metaTitle` ≤51 chars (the `%s - Servey` template adds 9; rendered ≤60).
- Descriptions 110-160 chars.

- Explicit `alternates.canonical` and explicit `openGraph.url` on every page.
- One `<h1>`, no skipped heading levels.
- Shared `ogImage` object (never the bare string, which drops width/height/alt).

Technical

- **Static generation everywhere.** Every route is Static/SSG; guard this.
- **Speculation Rules** - same-origin prefetch/prerender at `moderate` eagerness, excluding `/api/` and `/ingest/`.
- **Structured data:** one `@graph` (Organization → WebSite → SoftwareApplication with PreOrder Offers → FAQPage → ImageObject), plus per-page `Article` / `WebPage` / `BreadcrumbList` / `FAQPage`. Cross-reference by `@id`; never re-declare a node inline.
- **sizes on every responsive image**, matched to its real rendered width.
- **IndexNow** after every content deploy (`npm run indexnow`) - fastest route into Bing, and therefore into ChatGPT Search.
- `npm run seo:audit` before/after deploys - all sitemap URLs pass at last run (run it against a local build with `--base=http://localhost:PORT` to cover pages that are not on the live sitemap yet - the default base is production, so new posts are silently skipped on the run that matters most).
- `npm run seo:meta` (`scripts/meta-audit.mjs`) reads the rendered `.next` HTML and fails on duplicate titles, duplicate descriptions, and length violations. It catches what `seo:audit` cannot: collisions only visible when every page is compared against every other. It found `/headless-mac-mini` at 171 chars within a day of existing.
- **CI is wired** (`.github/workflows/checks.yml`, 28 Sep): `tsc`, `lint`, `build`, `seo:meta`, `seo:framing` on every push and PR to `main`. `seo:framing` is `continue-on-error` because of the backlog in §11 - tighten it to a hard gate once that clears, or the signal is worthless.
- **The AI files are generated, not written.** `scripts/build-ai-files.mjs` runs as `prebuild` and emits `public/llms.txt` (index), `llms-full.txt` (~174KB, every post in full) and `ai.txt` from `lib/blog.ts`, `lib/use-cases.ts` and `lib/content.ts`. This exists because the hand-written `llms.txt` had silently drifted: 12 posts missing and a retired overclaim still in it a month after it was removed everywhere else. **Never edit the files in `public/` - the next build overwrites them.**
- **Redirects live in `next.config.ts`**, before `rewrites()`. Next emits **308** for `permanent: true`, not 301; Google treats the two identically.

Answer engines (AEO / GEO)

- **Self-contained passages of ~130–170 words** that still make sense quoted in isolation.
 - **Question-shaped H2s** mirroring real queries.
 - **Direct answer in the first sentence** after each heading, then support.
 - **`robots.ts` explicitly allows AI crawlers by name** - GPTBot, OAI-SearchBot, ChatGPT-User, ClaudeBot, Claude-Web, Claude-SearchBot, PerplexityBot, Perplexity-User, Google-Extended, Applebot-Extended, Bingbot, Amazonbot, CCBot.
-

4. What we deliberately DO NOT do

This section is as important as the last one. Each of these is a considered decision with a reason.

Never - integrity

We don't	Why
Invent ratings, testimonials, review counts, user numbers, or <code>AggregateRating</code>	Survey is pre-launch with zero public users. <code>AggregateRating</code> without real reviews is a Google structured-data violation, and the rest is simply false.
Fabricate product screenshots	Every image is a real capture or an honestly-styled placeholder.
Publish a comparison that hides where a competitor wins	It is the entire credibility asset. A post that reads like an advert has failed at its job.
Imply Termius shipped malware	They were the victim of the ZuRu impersonation campaign. Saying otherwise is false and defamatory.
Solicit upvotes anywhere	Against Product Hunt / AlternativeTo rules; triggers rank penalties.
Add our own links to Wikipedia (dead-link building)	Verified 24 Aug: Wikipedia renders external links as <code>rel="nofollow"</code> sitewide, so they pass no ranking signal. Adding our own also violates WP:COI/WP:REFSPAM, where the failure mode is <code>servey.in</code> on the public Wikipedia spam blacklist. Wikipedia is downstream of independent press, not a substitute. Broken-link building off Wikipedia is fine and is our active tactic - see <code>docs/BACKLINKS.md</code> .
Buy links, use PBNs, or exchange links	Cheap short-term, fatal long-term.
Work a bulk "N places to get backlinks" list top to bottom	Triaged one of 75 on 26 Aug: 9 were usable now . The rest were editorial placements you must earn (Forbes, TechCrunch), revenue-gated directories we cannot enter pre-launch, review sites needing real customers, open-source-only listings, or content farms. Submitting to dozens of unrelated low-quality directories in a burst is the classic footprint of a link scheme, and on a domain with 2 referring domains it is the <i>only</i> pattern in our profile. Pick for relevance, submit steadily, and never post the same text everywhere without a canonical. See <code>docs/BACKLINKS.md</code> §6.
Keyword-stuff, cloak, or hide text	Trivially detected; the fragmentation fix is pillar pages, not stuffing.
Publish thin AI-generated filler to inflate page count	Google's QRG §4.6.5–6 targets exactly this. Page count is not the goal.

Not now - strategy

We don't	Why
Chase "recommend me a tool" AI queries	Verified 16 Aug: ChatGPT answered "tool to screen share to my Mac from my phone" with Chrome Remote Desktop, AnyDesk, TeamViewer. Correctly - Servey has nothing to download, so recommending it would be a worse answer. Structurally unwinnable until launch. Do not treat it as an SEO failure.
Chase "termius alternative" broadly	It would pull multi-host-SSH searchers who bounce, polluting waitlist quality pre-launch. The honest overlap is narrow: "run a terminal on your own Mac without SSH setup".
Compete on Termius's page axis (SSH client x platform)	They hold #1 there. We copied their <i>template technique</i> onto a different axis - task x device - where they do not rank at all.
Add more pages targeting "control Mac from iPhone"	That recreates the exact fragmentation <code>/control-mac-from-iphone</code> was built to fix. One pillar, many spokes.
Rewrite metadata sitewide to chase CTR	Still ruled out as a <i>bulk</i> exercise. But the original reasoning ("above par for position 16.4") used the misleading site average and is void - see §1. Targeted work on the specific page-one comparison URLs is now allowed, and only those.
Build doorway pages per query variant	The 10 phrasings get absorbed by one strong page, not 10 weak ones.
Do multilingual / translated SEO	Considered 24 Aug and deliberately held. 19 posts x 7 locales is 133 URLs on a domain with 2 referring domains already averaging position 16.3. Translations multiply pages, not authority, and scaled machine translation without human review is named in Google's spam policies. The tactic works for language-independent utility queries ("instagram photo downloads"); ours are niche Mac comparison queries whose product names, docs and buyers are all English. Revisit once referring domains are in double digits, then as a one-language, 4-post pilot chosen from GSC Countries data - never a 7-locale bulk translation.
Chase <code>airplayuiagent</code> traffic	Accidental, off-intent, and converts nobody.

Myths we do not act on

Belief	Reality
llms.txt improves AI citation	It does not. No major engine consumes it; this is evidence-based, not opinion. We keep ours because it costs nothing and is a tidy human-readable map - but it is not a citation lever and must not be cited as one.
FAQ schema wins SERP real estate	Retired 7 May 2026. See §7.
More pages = more traffic	39 queries across 26 pages says relevance is fine; authority is the constraint.
Content needs chunking / AI-specific keyword rewriting	Engines handle synonyms; this is wasted effort.

5. Cluster → page ownership map

Each intent has exactly one owner. Adding a second competitor for any row is how fragmentation comes back.

Intent cluster	Owning page	Supporting spokes
Control Mac from iPhone	/control-mac-from-iphone	blog/control-your-mac-from-iphone-ipad (the how-to half - see §2)
Control iPhone from Mac (inverse intent)	blog/control-iphone-from-mac	-
Terminal from iPhone	/terminal-on-iphone	blog/real-terminal-on-your-mac-from-iphone
Terminal from iPad	/terminal-on-ipad	-
Headless Mac mini / macOS headless mode	/headless-mac-mini	blog/headless-mac-mini-setup (setup), blog/control-a-headless-mac-mini-remotely (access)
Run local AI on Mac / mac ai agent	/remote-mac-for-ai-agents	blog/run-ai-agents-locally-on-your-mac , blog/run-ai-agents-on-your-mac-remotely , blog/stay-in-control-of-ai-agents-from-anywhere
Developer remote Mac access	/mac-for-developers	blog/who-is-servey-for-developers-home-labs
Mac home lab	/mac-home-lab	-
Cellular / strict networks	blog/access-your-mac-remotely-over-cellular	-
Screens vs Jump Desktop	blog/screens-vs-jump-desktop	blog/screens-jump-desktop-alternative-mac
Back to My Mac replacement	blog/what-replaced-back-to-my-mac	blog/does-mac-screen-sharing-work-over-the-internet
RustDesk vs AnyDesk	blog/rustdesk-vs-anydesk	-
Screens 5 alternatives	blog/screens-5-alternatives	-
Jump Desktop vs TeamViewer	blog/jump-desktop-vs-teamviewer	-
Jump Desktop vs RustDesk	blog/jump-desktop-vs-rustdesk	-
Termius / SSH overlap	blog/termius-alternative-mac-terminal	-
Category roundup	blog/best-remote-desktop-for-mac	-

6. Internal linking - a quirk that matters

`p` blocks render as plain text. `app/blog/[slug]/page.tsx` parses no markdown and no HTML, so **you cannot put an anchor in body copy**. The only real internal-link mechanism between posts is **exact-string keyword overlap** feeding `relatedPosts()`.

Engineering the `keywords` array *is* engineering the link graph. When adding a post, deliberately share keywords with the posts you want linking to it. The three newest comparison posts were given keyword overlap with `screens-vs-jump-desktop` specifically because it is the one page earning clicks, so its authority flows to them.

Use-case pages *do* have real anchors: the cross-link mesh is derived automatically from the `useCases` array order, so adding a page wires it in.

7. Structured-data deprecation register

Type	Status	Action
FAQPage	Rich results retired 7 May 2026. SC report + Rich Results Test support removed June 2026; API data August 2026.	Keep the markup. It remains a valid Schema.org type, Google confirms unused structured data does not affect Search, and answer engines still read it. Only the <i>expectation</i> changed.
HowTo	Removed June 2025	Do not add
SpecialAnnouncement , ClaimReview , VehicleListing , EstimatedSalary , LearningVideo , Course carousel, Book Actions	Removed June 2025	Do not add
QAPage	Retained	Available if ever useful
VideoObject	Active	Add when the YouTube demo exists
AggregateRating	Active	Forbidden here - no real reviews exist

8. Constraints that shape everything

1. **Pre-launch, waitlist-only, nothing downloadable.** This is the dominant constraint. It makes purchase-intent and "recommend a tool" queries unwinnable, and it blocks the highest-authority directories (§9).
 2. **Effectively no backlinks.** The real reason for position 15.8.
 3. **Almost no entity presence.** AI engines weight presence on YouTube, Reddit, LinkedIn and Wikipedia. We have essentially none - and **three of those four are free and available pre-launch.**
 4. **The brand name carries a search tax.** Bing responds to "servey" with *"Do you mean **survey**?"*. Google brand search works (131 impressions), so it is not fatal, but every search box and answer engine will fight the name. Worth knowing before blaming rankings for it.
 5. **hero-devices.png** **exposes personal filenames** and is live on the homepage. The owner's decision is to leave it. Not an SEO issue; recorded so nobody "fixes" it by accident.
-

9. Channels - available vs blocked pre-launch

Verified against each platform's own rules.

Channel	Status	Note
AlternativeTo	 Blocked	<i>"If the app is Coming Soon (announced only), in closed/private beta, invite-only, or in Early Access phase, we do not accept it."</i> Account is created and past the 7-day age requirement, so submission can go in on launch day.
Show HN	 Blocked	<i>"If your work isn't ready for users to try out, please don't do a Show HN."</i>
App Store listing	 At launch	#1 most-cited source in this category (77 citations)
YouTube demo	 Available	#2 most-cited (74 citations). The highest-leverage action available today, and still not done.

See `docs/DEMO-VIDEO.md` for what the existing footage can and cannot be used for, and the shot list for recording it properly. | **Reddit** (r/macapps, r/homelab) |  Available | **#3 most-cited** (65 citations). Genuine participation, never a launch post | | **BetaList / Peerlist / Uned** |  Available | Built specifically for unlaunched products | | **HN as a normal submission** |  Available | Show HN is barred; submitting a post as a regular link is fine | | **Email the waitlist** |  Available | Still never used | | **Product Hunt** | Done | Badge in footer. Never solicit upvotes |

10. Measurement

Tool	What it answers	How
<code>npm run seo:audit</code>	Do the on-page invariants still hold?	Titles ≤ 60 , descriptions 110–160, canonical self-reference, og:url agreement, one h1, no skipped headings, valid JSON-LD with resolving <code>@id</code> s, BreadcrumbList count == visible trail , all sitemap URLs 200. Exits non-zero - can gate a deploy.
Google Search Console	Impressions, position, query clusters	The cluster table in §2 is the useful view, not the top-line number
Bing / ChatGPT	Are we retrievable by answer engines?	Phrase search on Bing; test comparison queries, not "recommend me a tool"
PostHog	What happens after the click	<code>\$pageview</code> , <code>waitlist_submitted</code> → <code>waitlist_signup</code> (with <code>duplicate</code>), <code>waitlist_failed</code> , all carrying <code>source</code>
<code>npm run seo:meta</code>	Are any two pages wearing the same title or description?	Over the rendered <code>.next</code> HTML. Needs a build first. Exits non-zero.
<code>npm run seo:framing</code>	Is Servey buried, or laid on too thick?	Depth $\leq 15\%$, substance $\leq 35\%$, density 4.0–6.0/1k (<700-word posts to 9.0). Comparison and roundup posts are exempt from depth and substance by design.
CI (<code>.github/workflows/checks.yml</code>)	Did a push break any of the above?	tsc, lint, build, <code>seo:meta</code> , <code>seo:framing</code> on push/PR to <code>main</code>
<code>npm run indexnow</code>	Push new URLs to Bing	Run after every content deploy

Reading GSC correctly

- **"Discovered – currently not indexed"** → Google knows the URL, hasn't crawled it. Requesting indexing helps.
- **"Crawled – currently not indexed"** → Google crawled it and *chose* not to index. Requesting indexing will **not** help. That is an authority verdict, and the correct response is links and entity presence, not more pages.

11. What to do next, in priority order

1. **Record the YouTube demo.** Highest-leverage action available pre-launch, feeds YouTube + Reddit + the launch-day AlternativeTo submission, and is the only top-3 cited source reachable before shipping.
2. **Close the Bing crawl gap (new, 1 Sep).** The comparison posts sit in `Discovered but not crawled`, so ChatGPT Search - which reads Bing's index - cannot see the only cluster we rank page one for. Fix via **URL Inspection -> Request indexing**, one URL at a time. **Not** via Site Explorer, which only lists URLs Bing already crawled and therefore cannot reach these. **Status 1 Sep: 16 URLs submitted (see change log); sitemap cleaned and resubmitted by the owner - 1 known sitemap, 0 warnings, 34 URLs discovered, crawled 1 Sep.** What remains is waiting for Bing to actually crawl: re-inspect a sample in ~a week, starting with `splashtop-vs-jump-desktop` and `control-your-mac-from-iphone-ipad`. **5 Sep: done - 12 URLs submitted (4 new, 4 rewritten, 4 minor) and the sitemap resubmitted; quota 100 -> 88 confirms all 12 landed.** Two things still owed: the 1 Sep batch has now had 4 days, so re-inspect that sample; and GSC Request Indexing for the 8 tier-1 and tier-2 URLs is **the owner's to do** - Bing is done, Google is not.
3. **Add CI (new, 10 Sep). There is none -- no `.github` directory at all. Done 28 Sep.**
`.github/workflows/checks.yml` runs `tsc`, `lint`, `build`, `seo:meta` and `seo:framing` on every push and PR to `main`. It was worth doing because every guard this repo had was run by hand, which is why several regressions were caught by eye rather than by a gate: the overclaim surviving in `public/llms.txt`, the seven-card grid, a post shipping at 2.9 density. **What is still owed:** `seo:framing` runs `continue-on-error` because of the pre-existing failures in item 5 below, so it currently reports rather than blocks. Clear those, then drop the flag - a check that is allowed to fail teaches people to ignore it.
4. **Add in-app account deletion (new, 6 Sep) - App Store blocker.** Guideline 5.1.1(v) has required apps that support account creation to offer account deletion *inside the app* since 30 June 2022. Servey has none: there is no delete path in either app, and `firestore.rules` sets `allow delete: if false` on `users`. Temporarily disabling an account does not satisfy it - the data has to go. This is a rejection at submission, not a nice-to-have, and it is cheaper to build now than during review. The privacy policy currently documents deletion by emailing `hello@servey.in`, which is honest and legally sufficient but is **not** what Apple checks for.
5. **Fix the posts `seo:framing` flags (new, 5 Sep, still open).** The script surfaced pre-existing failures; **none were touched**, because rewriting a dozen posts was not what was asked. **This is now also what keeps CI honest**
 - `seo:framing` runs `continue-on-error` purely because of this backlog (item 3). As of 28 Sep the run is **17/28 passing**. Worth doing in priority order:
 - **`control-your-mac-from-iphone-ipad` (substance 70%) and `headless-mac-mini-setup`.** Both were explicitly "fixed" on 31 Aug for *depth* - 4% and 3% - and both still bury Servey. Same blind spot, same posts, which is the strongest evidence the new metric was needed. `headless-mac-mini-setup` absorbed the merged post on 28 Sep, so re-measure it before rewriting - the merge moved its substance figure.
 - **5 short posts at 69-79% depth:** `real-terminal-on-your-mac-from-iphone`, `access-your-mac-remotely-over-cellular`, `run-ai-agents-on-your-mac-remotely`, `run-ai-agents-locally-on-your-mac`, `stay-in-control-of-ai-agents-from-anywhere`. All 360-575 words, all naming Servey only in a closing section. (`control-a-headless-mac-mini-remotely` was on this list and no longer exists - it was merged away on 28 Sep.)

- **Density too high:** `who-is-servey-for-developers-home-labs` (16.6) and `termius-alternative-mac-terminal` (16.4) are Servey-centric by design, so check whether the *post* is right before changing the number. `splashtop-vs-jump-desktop` (6.7) and `chrome-remote-desktop-vs-jump-desktop` (6.5) are marginal.
6. ~~Capture the Mac host's SESSIONS screen.~~ **Done 5 Sep** - owner supplied it, cropped to the window bounds and shipped as `public/screenshots/terminal-sessions.png`. The persistent-sessions card now has a real capture, and no feature card ships a placeholder any more.
 7. **Reddit participation** in `r/macapps` and `r/homelab` - genuine, not promotional.
 8. ~~Email the waitlist.~~ **Ruled out by the owner** - no bulk or broadcast sends to existing waitlist members. Do not re-propose this.
 9. **BetaList / Peerlist / Unneed** submissions.
 10. **Add `VideoObject` schema** once the demo exists.
 11. Re-test AI citation on **comparison** queries in ~2 weeks.
 12. At launch: App Store listing, AlternativeTo submission, Show HN.

Explicitly not on this list: more landing pages, more metadata rewrites, more schema. On-site SEO has reached sharply diminishing returns. What is missing is that **nobody has ever seen Servey move.**

12. Change log

Date	Change
2026-09-29	Benchmark harness built, and the docs caught up with a month of tooling. Harness (<code>scripts/benchmark/</code>): <code>flash.html</code> and <code>pattern.html</code> as stimuli, <code>measure-latency.py</code> (frame-differencing a phone recording of Mac screen and client side by side) and <code>measure-sharpness.py</code> (SSIM of the rendered pattern against the source). SSIM self-tested before use - identical images return 1.0000, a 1.2px Gaussian blur returns 0.9616, so the metric discriminates at the scale that matters. The post draft lives at <code>docs/benchmark-post-draft.md</code> and is deliberately absent from <code>lib/blog.ts</code> : every figure is a literal <code><<FILL: ...>></code> , so it is structurally impossible to publish with invented numbers. Screens and Jump Desktop both have to be measured on the same network in the same sitting or the comparison is worthless. Docs audit found four gaps and closed them: both files still said 38 sitemap URLs against an actual 40 (28 posts, 7 use cases, <code>/mac</code> , <code>/</code> , <code>/blog</code> , <code>/privacy</code> , <code>/terms</code>) and CONTEXT still said 23 posts; and three pieces of tooling existed in no document at all - <code>npm run seo:meta</code> (<code>scripts/meta-audit.mjs</code>), <code>.github/workflows/checks.yml</code> , and <code>scripts/build-ai-files.mjs</code> wired to <code>prebuild</code> with its generated <code>llms-full.txt</code> / <code>ai.txt</code> . The last one is the dangerous omission: <code>public/llms.txt</code> is now generated , and anyone editing it by hand would have their work silently overwritten by the next build - which is the inverse of the drift that caused the generator to be written. Lesson, same shape as the <code>public/</code> sweep miss on 5 Sep : tooling added under time pressure does not document itself, and a doc that is 80% current is trusted like one that is 100% current. Audit the counts in the docs whenever the page count changes.
2026-09-28	Near-duplicate merged, and an image pass. GSC 3 months to 25 Sep: 183 clicks, 10.1K impressions, position 9.0 (87 / 4.28K / 12.7 on 1 Sep). The pattern worth acting on : comparison and buying-decision pages convert at 3-6% while question-answering pages convert near zero at the <i>same</i> positions - <code>does-mac-screen-sharing-work-over-the-internet</code> is now the site's highest-impression page (882) at 0.8% from position 8.1, and a title rewrite on 10 Sep moved it barely at all. Most likely those questions are answered in the results page itself, so no metadata work reaches them. Treat pure Q&A as a low-yield format here and keep writing comparisons; the two newest posts (iPhone Duo 4.9%, new Mac mini 6.1%) are the best performers by rate and both came off a real product event. Merged <code>control-a-headless-mac-mini-remotely</code> (328 impressions, 1.5%, position 9.8, ~360 words) into <code>headless-mac-mini-setup</code> (569, 1.6%, 6.3): the guide absorbed its one distinct argument (you need a screen as well as a shell, because the blocker is sometimes a dialog box), two of its FAQs rewritten in the guide's neutral voice, and its one unique keyword. 301 redirect in <code>next.config.ts</code> , and the <code>headless-mac-mini</code> use case's <code>relatedSlug</code> repointed. 29 posts -> 28. Images : every screenshot converted to WebP (6.5MB -> 424KB; the hero alone was a 3.9MB 2560x1440 PNG, now 1920x1080 at 171KB with the registry dimensions updated so object-cover does not crop), two unreferenced files deleted including the one showing an email address, an image sitemap added (homepage lists every ready shot, <code>/mac</code> the host app, each post only what it embeds), and the connection diagram's alt stopped saying "HEVC ... WebRTC". Delivered bytes barely change - <code>next/image</code> was already serving WebP/AVIF - so the wins are image-search eligibility and 6MB off the repo.
2026-09-10	GSC pulled properly for the first time since 1 Sep: two findings that change what to work on. (1) The <code>survey</code> query is not brand demand - it is people misspelling "survey". 347 impressions, position 4.9 , and 2.9% CTR . A brand term at position 5 converts at 15-30%; 2.9% means 337 people saw us and did not want us. It is 8% of all site impressions , it is the single largest query, and it is what drags the headline "2% average CTR" down. Do not optimise for it, do not read the sitewide CTR as a health metric, and do not treat rising impressions as rising interest until this is subtracted. (2) The real opportunity is CTR on page-one comparison pages, and it is large. Per-page, 3 months: <code>does-mac-screen-sharing-work-over-the-internet</code> 394 impressions (2nd highest on the site), position 9.7, 1.0% CTR ; <code>screens-vs-jump-desktop</code> 365 / 6.2 / 3.8%; <code>jump-desktop-vs-rustdesk</code> 284 / 7.2 / 3.9%; <code>screens-5-alternatives</code> 241 / 7.0 / 2.9%; <code>splashtop-vs-jump-desktop</code> 200 / 6.9 / 2.0%. These

Date	Change
	already rank page one. Positions 6-10 should return 5-8%; we are getting 1-4%. Lifting the five to a normal band is roughly +60 clicks a quarter against a current total of 87 - a near-doubling with no new ranking and no new posts. This is the evidence §4 was waiting for: the CTR ban was justified by a blended 16.5 average position that mixed two unrelated populations, and per-page data at known positions is the thing that was missing. Acted on it: rewrote <code>metaTitle</code> and <code>description</code> for all five, keeping the head keyword in every title so rankings are not put at risk, and replacing "which should you pick" style restatements of the query with the actual deciding axis (bought once vs subscription; is free good enough; the four fixes). (3) The core cluster is still stranded: <code>remote control iphone from mac</code> position 42.5, <code>remote access mac from iphone</code> 38.7, both 0 clicks on 57 impressions. Page four earns nothing, and no metadata work reaches it - that needs links, which is unchanged. (4) Trend is genuinely up: 4.28K impressions over 3 months, 328 on 7 Sep alone, against a near-flat July. Method note: GSC defaults to clicks-only columns; enable Average CTR and Average position or the per-page diagnosis is impossible. Sorting by impressions rather than clicks is what surfaced both findings.
2026-09-06	Privacy Policy and Terms rewritten for real; both had opened by calling themselves placeholders. Every factual claim in the privacy policy was read out of the app repo, not assumed - the same discipline used on the feature claims. Findings worth keeping: the apps carry no analytics SDK and no crash reporter at all (Firebase Auth + Firestore, Supabase, Google Sign-In, StoreKit, and nothing else), billing is Apple-only so card data never reaches us, and the WebRTC signaling exchange contains IP addresses - undisclosed anywhere until now, and now its own subsection, stated as ephemeral because the code deletes the session doc. Supabase holds session counts and total seconds, so the policy says counts rather than "usage data", which invites a reader to assume content. The <code>users</code> table has a <code>phone</code> column nothing writes to, so we do not claim to collect phone numbers. Structure follows <code>macky.dev/privacy</code> where it fits plus what DPDP 2023 requires of a notice (purposes, categories, retention, withdrawal, rights procedure, grievance route, 90-day resolution) and GDPR/UK GDPR rights. Ownership and discontinuation are in both documents at the owner's request: privacy covers what happens to data if the service ends, terms carry the enforceable version. Terms also gained acceptable use, Apple-handled refunds, limitation of liability and Indian governing law - it had none. Blocker found, not fixed: the app has no account deletion (<code>allow delete: if false</code> on <code>users</code>, and no delete path in either app). App Store Review Guideline 5.1.1(v) has required in-app account deletion since 30 June 2022, so this is a submission rejection waiting to happen; the policy currently documents deletion by email, which is honest but is not what Apple asks for. See §11.
2026-09-05	Bing: 12 URLs submitted and sitemap resubmitted. Worked out what actually changed rather than resubmitting everything - 4 new (the new posts), 4 substantially rewritten (<code>/</code>, <code>/blog</code>, <code>/privacy</code>, <code>/remote-mac-for-ai-agents</code>), 4 one-sentence edits (the 3 posts carrying the retired overclaim, plus <code>control-iphone-from-mac</code>'s description fix). Method note: diffing <code>lib/blog.ts</code> naively reported 14 changed posts, all false - prettier had reflowed the file. The reliable check is to grep for the <i>specific replacement strings</i> and map each hit back to its slug, or compare rendered HTML; never trust a source diff after a formatter has run. All 4 new posts were Discovered but not crawled, discovered 05 Sept - that is IndexNow working (it got them into Bing's queue within hours) and simultaneously <i>not</i> working (queued is not crawled), which is the same split logged on 1 Sep. <code>/</code> was Indexed successfully but holding the pre-rebuild page - 6 features, 2 plans, no free tier, and the overclaim - so the recrawl request matters more there than anywhere else. Quota went 100 -> 88, exactly 12, which is the check to use; the success toast is easy to miss when batching. Sitemap re-submitted: last submit 9/5, status Processing, still 1 known sitemap / 0 errors / 0 warnings, and it will move 34 -> 38 once processed. GSC left to the owner. Advice given: spend GSC's ~10-12/day Request Indexing quota on the 8 in tiers 1-2 only and skip the 4 one-sentence edits, which Google will pick up unaided.
2026-09-05	Owner: the new posts over-served the free how-to and buried Servey. Both true; the depth metric could not see the second one. Depth measures the first <i>mention</i>, which a table cell or the lede satisfies - all four posts read 4.9-9.0% while the first paragraph that actually says what Servey does sat at 49%,

Date	Change
	51%, 65% and 69%. A reader on a phone scrolled past ~700 words of free tutorial first. This is the 31 Aug regression in a new costume, and the rule written to stop it was satisfiable by a token, so a third metric now exists: substance (first ≥ 25-word paragraph or <code><h2></code> naming Servey, lede excluded) with a $\leq 35\%$ ceiling, in <code>scripts/servey-depth.mjs</code> / <code>npm run seo:framing</code>. Fixed structurally, per the 26 Aug precedent: a short-answer paragraph after each opening table giving the real answer and naming Servey with a "the app we make" disclosure, plus naming Servey in the privacy-screen table's approach row (the only one of the four tables that did not). Cut, because the owner was right it was surplus and it was what pushed Servey down: <code>pmset list 5→3</code> commands, the dummy-HDMI and heat/battery/FileVault paragraphs, the screen-vs-tmux paragraph, the shopping advice, the remediation walkthrough, <code>ls of</code> and two log commands - every one of them duplicated a table row or its own FAQ answer verbatim. Then removed 7 repeated Servey mentions, because trimming words while holding mentions took the closed-lid post to 9.3/1k, which reads as an advert on exactly the pages the broken-link outreach points at. Final: depth 4.9-9.0%, substance 18-32%, density 5.8-6.2. <code>readingMinutes</code> corrected (closed-lid claimed 7 min for 971 words against a ~235 wpm site convention). The script found pre-existing failures I have not touched - see §11.
2026-09-05	Site brought up to date with the app; 34 -> 38 URLs. The site described the product as of roughly mid-July while four things had shipped since, so the copy work came before the content work. Added: a 7th feature card for persistent sessions (named tmux sessions on the Mac that outlive the app - the biggest change to the product and the one that converts the AI-agent persona already written into the page, which until now implied you had to sit and watch); a Free column in pricing (5-minute sessions, 5 a day, no card - the entire top of the funnel, previously invisible); the second lock in the privacy card and policy (master password set on the Mac, per-device approval, revocation - the page said only "device registration & approval"); and the activity log in the privacy policy. Corrected: "video hardly touches our servers" was an overclaim and appeared in 8 places. True on LAN and for direct P2P, false on a strict NAT where the stream relays - which the Networking card admitted two sections earlier. Replaced everywhere with "P2P first; our own relay if your network insists", which is <i>stronger</i> now the relay is ours (<code>turn.servey.in</code>) rather than a vendor's. Also: launch FAQ no longer reads as less certain than the pricing section three scrolls above it; macOS 15.3 / iOS 18.5 stated plainly (narrow, excludes visitors, better learned before joining a waitlist than after); Terminal plan bullets rewritten (oldest copy on the page, and it undersold what the tier became). Four posts, each aimed at existing demand rather than at a feature name: <code>keep-terminal-session-running-after-disconnect</code> (SIGHUP, tmux, and why it is harder from a phone - no Control key for the prefix), <code>blank-mac-screen-during-remote-access</code> (why vendor privacy screens keep breaking on macOS - kernel/DriverKit extensions - and why the gamma-after-capture approach avoids every overlay failure mode; targets <code>teamviewer black screen mac</code>), <code>macbook-closed-lid-remote-access</code> (clamshell, dummy HDMI, what <code>pmset disablesleep</code> really changes), <code>who-connected-to-my-mac-remotely</code> (<code>last</code> vs <code>log show</code> , and the honest point that the unified log rolls over in hours). Depth 3.6-6.2%, density 5.4-6.0/1k, all measured before committing per the 1 Sep rule. Every feature claim was verified against the app repo at 894c87f , not against a changelog - free-tier numbers read out of <code>SessionLimitManager</code> , deployment targets out of the <code>pbxproj</code>, relay host out of the ICE config. The audit that started this is an artifact the owner commissioned; treating its claims as requirements rather than facts was the right instinct, because two of them needed the source to confirm. Fixed <code>control-iphone-from-mac</code> at 161 chars, the one failing audit check. Local audit 38/38, zero duplicate titles or descriptions across 40 rendered pages.
2026-09-01	Bing sitemap cleaned by the owner. The 7 bogus rows (page URLs submitted as sitemaps on 16 Jul) deleted and <code>sitemap.xml</code> resubmitted. Bing now reports 1 known sitemap, 0 errors, 0 warnings, 34 URLs discovered, last crawl 1 Sep - matching the live sitemap exactly and replacing the stale 19 Aug read at 29 URLs. Both remaining Bing inputs (sitemap freshness and per-URL crawl requests) are now clean; the only open Bing item is waiting for the crawl itself.

Date	Change
2026-09-01	Bing crawl gap worked through: 16 URLs submitted for indexing. Audited all 18 blog posts individually through URL Inspection and submitted every one that was either uncrawled or holding stale content. 12 were Discovered but not crawled - Bing knew them and had never fetched them, including 5 of the 6 comparison posts that rank page one on Google. 4 more were indexed but stale and got a recrawl request, most importantly <code>control-your-mac-from-iphone-ipad</code>, which Bing was holding as the pre-rebuild pitch version. Only <code>screens-vs-jump-desktop</code> and <code>best-remote-desktop-for-mac</code> were left alone - indexed and unchanged. Bing's daily quota is 100, not 10 (owner was right); it went 99 -> 83, which independently confirms all 16 submissions landed and is the check to use, since the success toast is easy to miss in a batch. Expect crawling over the following days; re-inspect a sample rather than assuming.
2026-09-01	Correction to the same-day Bing entry below: Site Explorer undercounts, and "9 of 34 indexed" was wrong. Checking individual URLs through URL Inspection shows <code>screens-vs-jump-desktop</code> "Indexed successfully" despite being absent from Site Explorer under both the <code>Indexed URLs</code> and <code>All URLs</code> filters. So the claim "not one comparison post is in Bing's index" was false, and it was drawn from a report that is not the index. Rule: Site Explorer is a lagging partial report; use URL Inspection for any per-URL question, and never conclude a URL is missing from Site Explorer alone. The real status of the uncrawled pages is Discovered but not crawled - Bing knows the URL and has not fetched it (<code>splashtop-vs-jump-desktop</code> discovered 19 Aug, <code>jump-desktop-vs-rustdesk</code> and <code>headless-mac-mini-setup</code> 31 Aug). This is not Google's <code>Crawled - currently not indexed</code>; it carries no quality verdict, it is a crawl queue, and Request indexing is the right lever. Two process notes: the IndexNow push earlier today did not clear these, so per-URL requests are needed on top of it; and the URL Inspection result panel does not refresh when you type a new URL and press Enter - it keeps showing the previous URL's verdict until you click Inspect, which briefly produced a false "Indexed successfully" for <code>headless-mac-mini-setup</code> here. Confirm the <code>urlToInspect</code> query parameter matches before believing a result.
2026-09-01	Bing Webmaster Tools reviewed. Its two "duplicate content" warnings are false positives; the real finding is underneath them. Both "too many pages with identical titles/meta descriptions" resolve to one cause: Bing indexes <code>https://servey.in/?ref=producthunt</code> as a page separate from <code>https://servey.in/</code>. Same page, tracking param on an inbound link nobody on our side generates (<code>grep</code> for <code>ref=producthunt</code> across the repo: no hits; not in the sitemap). Both URLs serve <code><link rel="canonical" href="https://servey.in"/></code>, verified live - Bing's SEO linter simply does not consult canonical before running this report. No code change; do not "fix" this by rewriting metadata or disallowing the param, which would break the canonical consolidation that is already working. Verified independently that the site has zero genuine duplicates: all 36 prerendered pages have unique <code><title></code> and <code><meta name="description"></code>, every title <= 60 chars, every description within 120-165 (script: build, then diff extracted tags). The real finding: Bing has indexed 9 URLs of 34 - see §1. Also confirmed the 7 bogus "sitemap" rows are still present (page URLs submitted as sitemaps on 16 Jul, all Warning / 0 URLs discovered); they are inert, not harmful, but they hide the one real row. IndexNow re-push of all 34 URLs sent and accepted.
2026-09-01	Audited against the eight durable SEO strategies; two real defects found and fixed, one standing diagnosis corrected. (1) Per-query position enabled for the first time - the site is bimodal, comparison cluster at 2.3-7.9 (page one) and core product cluster at 40-44.6 (page four). The 16.5 average we have reasoned from since 14 Aug is a midpoint between two unrelated populations. §1 core diagnosis rewritten; the §4 ban on CTR/metadata work is partially lifted, because it was justified with that bad average. (2) Internal links: 2 true orphans. Simulating <code>relatedPosts()</code> rather than counting raw keyword overlap showed <code>screens-vs-jump-desktop</code> - our best non-brand page, position 7.0 - receiving zero inbound internal links, because a top-3-by-overlap rule lets hubs crowd out mid-tier pages. Raising the count to 4 or 5 does not fix it and concentrates hubs further (max inbound 7 -> 10), so fixed at keyword level: <code>Screens vs Jump Desktop</code> added to <code>screens-5-alternatives</code> (30 Screens / 12 Jump Desktop

Date	Change
	mentions) and connect to Mac behind CGNAT to headless-mac-mini-setup (has a full section on CGNAT). All 23 posts now have >= 1 inbound; min 1, max 7. (3) Index bloat: none. 28 indexed, "Crawled - not indexed" 0. (4) Keyword cannibalization: checked and not present - keywords drives relatedPosts() and inert <meta keywords> , so 9 posts sharing best remote desktop for Mac is a hub, not competition. Real cannibalization would be two pages targeting one query in title/H1, and the /control-mac-from-iphone vs blog/control-your-mac-from-iphone-ipad pair is deliberately split commercial/informational.
2026-09-01	Survey framing regression on the 31 Aug posts, caught by the owner. The rule in §3 (first mention <=15% depth, density 4.0-6.0/1k) was written on 26 Aug in response to this exact complaint, and the two new posts broke it: control-iphone-from-mac 77% depth / 3.2 per 1k, headless-mac-mini-setup 76% / 1.8 - less than half the density floor. The specific failure the owner quoted was the phrase pattern "a VPN or an app that handles the networking, such as Survey": Survey demoted to a trailing example of a category, listed after the alternative, in a subordinate clause. Cause worth remembering: the §2 finding that day was that the old core page failed because it pitched instead of answering, so I over-applied "lead with the free built-in answer" and let it suppress Survey everywhere, including in FAQ answers where it does not apply at all. Leading with the honest free option is a rule about the opening frame of an informational section, not a licence to bury the product. Fixed to 6% / 5.2, 4% / 4.6, 3% / 5.2, all passing. Every honesty caveat verified intact afterwards (Survey does not control iPhones; Tailscale genuinely good; Apple-only; cannot be self-hosted; Screen Sharing free and already installed). Check depth and density before committing a post, not after.
2026-08-31	Link-graph fix on arrival. headless-mac-mini-setup shipped with 1 inbound post, the same failure mode logged on 23 Aug. Added headless Mac Mini remote control as a 6th keyword to real-terminal-on-your-mac-from-iphone (4 headless mentions, genuinely about driving one from a shell) and Mac Mini home server to who-is-survey-for-developers-home-labs (4 Mac mini mentions). Inbound 1 -> 3. control-iphone-from-mac arrived with 3 and needed nothing. Reminder for future posts: check inbound overlap before committing, not after - keyword overlap is the entire internal link graph (§6).
2026-08-31	GSC query count went 39 -> 162, which exposed the real problem. Pulled the full query list rather than the top 10 and clustered it (§2). The control my Mac from my iPhone family - Survey's exact job - was ~120 impressions across 55 phrasings, converting zero, the largest cluster on the site. Cause was legible on the page: control-your-mac-from-iphone-ipad was a pitch, not a guide, and never named System Settings, Screen Sharing or VNC while the queries were overwhelmingly how to ... Rebuilt it as a genuine how-to that leads with Apple's free built-in route, then places Survey against the away-from-home case it does not solve - the same shape that made does-mac-screen-sharing-work-over-the-internet convert. Added a 4-row approaches table, 7 FAQs replacing 3 pitch-only ones, metaTitle , and 7 keywords covering the actual phrasings. This page is the informational half of the pair; /control-mac-from-iphone stays the commercial half, so the rewrite reduces overlap rather than creating it.
2026-08-31	Two posts, 32 -> 34 URLs, both chosen from zero-click clusters rather than by adding more competitor comparisons. control-iphone-from-mac serves inverse intent - ~46 impressions across 13 phrasings of people wanting their iPhone from their Mac, which we had no page for and which was diluting the core page. Answered honestly with iPhone Mirroring (requirements verified against Apple support 120421) including the limitation that decides the question: the phone must be locked and near the Mac, so it is not remote access, and iOS permits no third-party equivalent. States plainly that Survey does not control iPhones before pivoting to the direction that does work. headless-mac-mini-setup serves ~32 impressions of setup intent (headless mac mini setup , macos headless mode) that the existing access-oriented post does not answer; real commands, and the FileVault-vs-auto-login trade-off stated as a trade-off rather than a recommendation. Link-worthy for the r/homelab outreach in BACKLINKS §2.

Date	Change
2026-08-31	Survey row highlighting. The green-border treatment added 27 Aug only detected a Survey <i>column</i> . Tables comparing <i>approaches</i> (rather than products) put Survey in a row, so the rebuilt core post got no highlight. Extended <code>app/blog/[slug]/page.tsx</code> to detect either orientation.
2026-08-31	Not written: a screens 5 cost post (~17 impressions, pure commercial intent, uncovered). Edovia does not publish pricing on the product page, and shipping figures we cannot verify - or vague copy that ranks for nothing - both fail. Revisit only with a verifiable source.
2026-08-26	Survey raised in the two Apple-tools posts. First-mention depth was 63% (<code>what-replaced-back-to-my-mac</code>) and 56% (<code>does-mac-screen-sharing-work-over-the-internet</code>), against 3% on <code>jump-desktop-vs-rustdesk</code> and 12% on <code>rustdesk-vs-anydesk</code> . The comparison posts score well because they lead with a table carrying a Survey column; the informational posts had tables with no Survey in them. Fixed structurally rather than by adding copy: the Back to My Mac jobs table gained a Survey column plus a terminal row (and an honest No on mounting the disk in the Finder), and the Screen Sharing approaches table - whose shape is approach x cost, so a Survey column would be a category error - instead names Screens, Jump Desktop, RustDesk and Survey inside the row they all belong to. Both posts now open the short-answer section by naming all four routes with a <code>which is the app we make</code> disclosure. 63% -> 14% and 56% -> 14% ; density 5.9 and 6.6 per 1k. The disclosure and the honest No are load-bearing here, not decoration: these are the pages the AppleInsider / MacRumors / 9to5Mac broken-link outreach points at.
2026-08-26	Two posts, 30 -> 32 URLs. <code>what-replaced-back-to-my-mac</code> exists to serve the broken-link campaign: AppleInsider, MacRumors and 9to5Mac all cite <code>support.apple.com/en-us/HT208922</code> , which still 404s (re-verified today, along with the other three dead Apple URLs). Our only replacement resource was <code>does-mac-screen-sharing-work-over-the-internet</code> , which covers Back to My Mac in a single section - Back to My Mac appeared 4 times sitewide, all inside that one post. An exact-match page makes the swap an easier yes for an editor. <code>rustdesk-vs-anydesk</code> feeds the only cluster earning non-brand clicks; RustDesk positions itself against AnyDesk explicitly, and the open-source axis already converts via <code>jump-desktop-vs-rustdesk</code> . Link graph engineered rather than incidental: <code>Back to My Mac replacement</code> added as a 6th keyword to the Screen Sharing post (it genuinely has a section and an FAQ on it), giving inbound counts of 6 and 3 on arrival against the 1 the previous post shipped with. Includes the AnyDesk 2024 production-systems incident, stated with its disclosure and cert rotation, because the comparison axis is literally auditable-vs-trust-the-vendor and omitting it would be dishonest. seo:audit 32/32.
2026-08-23	<code>does-mac-screen-sharing-work-over-the-internet</code> - first post in the Apple-tools cluster. Apple Remote Desktop appeared 0 times anywhere on the site; Screen Sharing and VNC appeared only as passing mentions inside comparison posts, with one keyword entry sitewide (<code>VNC alternative Mac</code>). No page targeted Apple's built-in tooling. Targets the Screen Sharing variant of the query (higher volume than the ARD variant Macky targets) and covers ARD and Remote Management inside it. Link graph: shipped with 1 inbound, so <code>macOS Screen Sharing</code> was added as a 6th keyword to the three posts that genuinely discuss it (<code>chrome-remote-desktop-vs-jump-desktop</code> 7 mentions, <code>best-remote-desktop-for-mac</code> 6, <code>screens-5-alternatives</code> 5) - inbound 1 -> 2, outbound now 3 topically-matched posts. 29 -> 30 URLs.
2026-08-23	Comparison posts re-pitched. The 6 comparison/roundup posts read as neutral review-site content: <code>jump-desktop-vs-teamviewer</code> had 2 Survey mentions in 1,526 words (1.3/1k), <code>anydesk-vs-teamviewer</code> 2 in 1,298. Their <code>Where Survey fits</code> sections actively deflected ("it does not compete with either of these", "pick one of the two above") and all 5 <code>Bottom line</code> sections had no Survey at all. Added a Survey column to all 5 head-to-head tables (each already ended with a `Real terminal

Date	Change
2026-08-19	Internal-linking fix: the 7 use-case pages were orphaned - they cross-linked to each other but nothing linked in, so the homepage emitted only 4 internal links (/ , /blog , /privacy , /terms). GSC counted 33 internal links sitewide with /privacy , /terms , /blog as top targets, i.e. the footer was the only link graph Google saw. Added a footer Use cases nav (homepage 4 -> 11 links) and a Servey for this block on posts inverting relatedSlug .
2026-08-19	3 comparison posts (splashtop-vs-jump-desktop , anydesk-vs-teamviewer , chrome-remote-desktop-vs-jump-desktop), keyword-wired to screens-vs-jump-desktop . 26 -> 29 URLs.
2026-08-19	Em dashes replaced with – across code comments and docs (owner preference). None were in rendered copy. source-material/ left untouched as owner input.
2026-08-16	3 comparison posts (screens-5-alternatives , jump-desktop-vs-teamviewer , jump-desktop-vs-rustdesk); /control-mac-from-iphone pillar; headless + local-AI pages retargeted; Speculation Rules; image sizes fix; ImageObject JSON-LD; og:image dimensions on all pages; /blog twitter card fix; Article publisher by @id ; scripts/seo-audit.mjs . 22 → 26 URLs. Sitemap re-read by Google: 26 discovered.
2026-08-16	PostHog analytics (lazy-loaded, first-party proxied, DNT honoured, no person profiles).
2026-08-14	6 use-case landing pages; Termius comparison; table block type; breadcrumbs; RSS feed; 404 page; light-mode contrast and reduced-motion fixes. 15 → 22 URLs.
Earlier	JSON-LD @graph ; IndexNow; AI-crawler allowlist; llms.txt ; comparison and roundup posts.